

Stat 443 Assignment 1: Exploratory Data Analysis

Caden Hewlett

January 30, 2024

Task 1: Analyzing Usual Hours Worked in Canada

a) Part (a)

Read in the data and create a time-series object. Plot the series and comment on any features of the data that you observe. In particular address the following points:

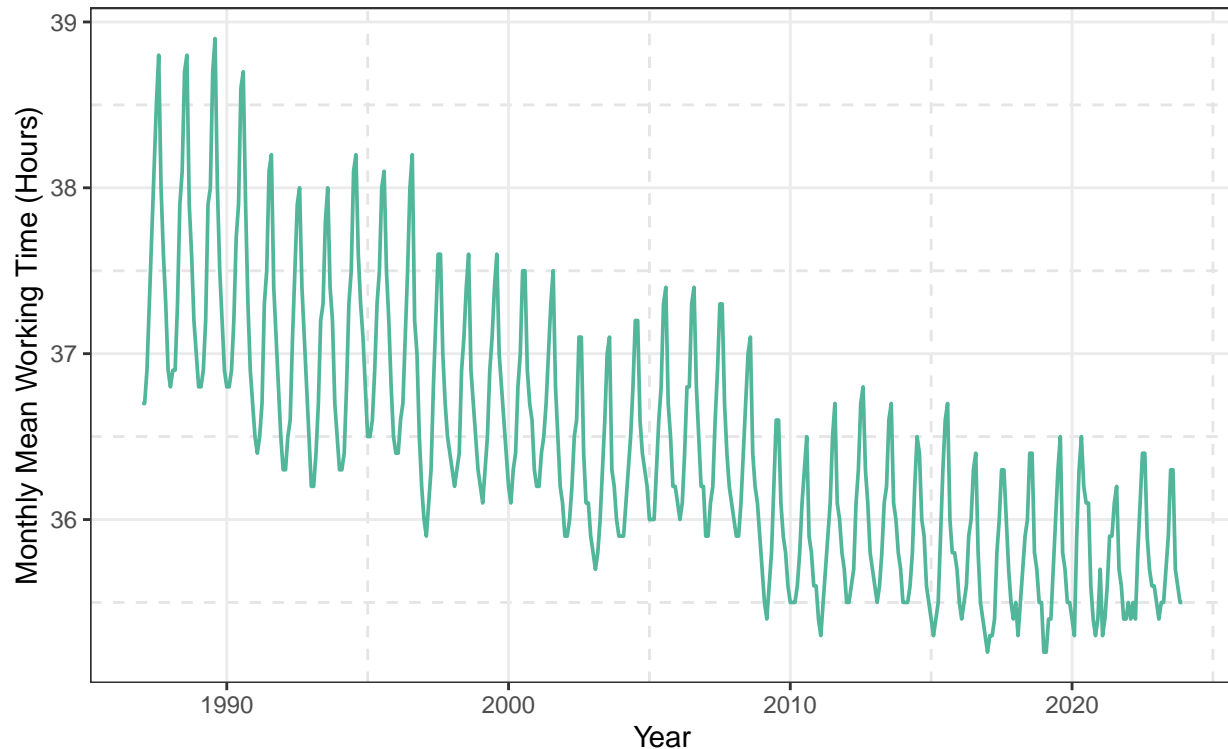
```
df <- read.csv("usual_hours_worked_ca.csv")

jobseries = ts(data = df$Hours, start = c(1987, 1), frequency = 12)

p1data = fortify.zoo(jobseries)
p1 <- ggplot(p1data, aes(x = Index, y = jobseries)) +
  geom_line(color = "#52b69a", linewidth = 0.65) +
  labs(
    title = "Monthly Average of Usual Hours Worked in Canada",
    subtitle = "Across all Industries from January 1987 to December 2023",
    y = "Monthly Mean Working Time (Hours)",
    x = "Year"
  ) + theme_bw() +
  theme(panel.grid.minor = element_line(
    color = "grey90",
    linetype = "dashed",
    linewidth = 0.5
  ))
print(p1)
```

Monthly Average of Usual Hours Worked in Canada

Across all Industries from January 1987 to December 2023



- Does the series have a trend?

Yes. It seems that there is a downward (negative) trend to the data, with the mean monthly average hours worked decreasing as a function of time. We would anticipate $m_t < 0$.

- Is there seasonal variation and if so would an additive or multiplicative model be suitable? Explain your reasoning.

Yes. There appears to be seasonal variation. Visually, we see this as a sinusoidal pattern to the time series. This seasonality could be caused by, for example, months during a given year.

Further, we would anticipate a multiplicative model, i.e. $\{X_t\} = m_t s_t Z_t$ for $Z_t \sim \text{WN}(0, \sigma^2)$ where $\sigma \in \mathbb{R}^+$. Visually, we can notice this by the changing amplitude of the seasonal periods over time.

To give an example of this, consider the following noise-less toy examples of Additive ($X_t = m_t + s_t$) Model and Multiplicative Model ($X_t = m_t s_t$)

```
st = 2*seq(from = 0, to = 8*pi, length.out = 1000)
mt = seq(from = 10, to = 0, length.out = 1000)
par(mfrow = c(1,2))
plot( mt + sin( st ), type = 'l',
      ylab = "Value of Xt", xlab = "Time (t)",
      main = "Additive")
plot( mt*sin(st), type = 'l',
      ylab = "Value of Xt", xlab = "Time (t)",
      main = "Multiplicative")
```



We see in the additive model, that the seasonal “amplitude” (i.e. the height of each peak/trough) does not change as a function of t , whereas in the multiplicative model the amplitude is changing due to the product with m_t . This is a key delineation of additive vs. multiplicative models. Hence, from our observation of the time series of the data, it is safe to assume that it is likely that the true $\{X_t\}$ takes a multiplicative model.

- Is the series stationary? Justify referring to the definition of a weakly stationary stochastic process.

This series is non-stationary. We can confirm this by the first property of a weakly stationary stochastic process, that $\exists \mu \in \mathbb{R}$ s.t. $\forall t \in \mathbb{N} \cup \{0\}, \mathbb{E}(X_t) = \mu$. By the presence of both seasonality and trend, we know that there cannot exist a real constant μ such that for all discrete time the expected value of the stochastic process is constant μ . This is because, by definition, m_t and s_t are functions of time. Therefore, $\mathbb{E}(t) = f(t)$ for some real-valued function f , and, since $\mathbb{E}(X_t)$ is a function of time, it cannot concurrently be some real-valued constant μ .

In more formal terms, for m_t and s_t being the trend and seasonal components of X_t respectively:

$$((\exists m_t \in \mathbb{R}) \vee (\exists s_t \in \mathbb{R})) \implies \nexists \mu \in \mathbb{R} \text{ s.t. } \forall t \in \mathbb{N} \cup \{0\}, \mathbb{E}(X_t) = \mu$$

Thus, the existence of either s_t or m_t denies the existence of μ . So, we know by the first property of weakly stochastic processes that this time series is non-stationary.

b) Create training and test datasets.

Part 1: The training dataset should include all observations up to and including December 2021; this dataset will be used to fit (“train”) the model. The test dataset should include all observations from January 2022

to December 2023; this dataset will be used to assess forecast accuracy. You can use the command `window()` on a `ts` object to split the data.

We'll start by splitting the data into train and test, then verifying our work. The verification process involves assuring that the sum of train and test is equal to the size of the series and also equal to the number of rows in the original data.

```
train <- window(jobseries,
  # starting at the beginning of 1987
  start = 1987,
  # ending at the end of 2021
  end = c(2021, 12),
  # monthly
  frequency = 12)

# get test data
test <- window(jobseries,
  start = c(2022, 1),
  end = c(2023, 12),
  frequency = 12)

# verify we've done things correctly
all.equal(length(test) + length(train),
  length(jobseries),
  nrow(df))
```

```
## [1] TRUE
```

Part 2 Using a suitable decomposition model and the loess method (R function `stl()`) decompose the training series into trend, seasonal, and error components. Plot the resulting decomposition.

We concluded in the previous question that this is likely to be a multiplicative rather than additive model. Hence, the best way to extract each individual component is to take the (natural) logarithm of the multiplicative model before decomposing the series into components. In other words, we will let:

$$\log(X_t) = \log(m_t s_t Z_t) = \log(m_t) + \log(s_t) + \log(Z_t)$$

So that we can consider each component separately. We will use `s.window = "periodic"`.

```
to_additive = log(train)
loess = stl(to_additive, s.window = "periodic")

seasonal = loess$time.series[, 1]
trend = loess$time.series[, 2]
noise = loess$time.series[, 3]
```

Then, we can plot the three components of the decomposition.

```
sp2 <- ggplot(data = fortify.zoo(seasonal), aes(x = Index, y = seasonal)) +
  geom_line(color = "#0077b6", linewidth = 0.5) +
  theme_bw() +
  labs(
    title = "Seasonal Component of Multiplicative Model",
    y = "log(st)",
    x = NULL
```

```

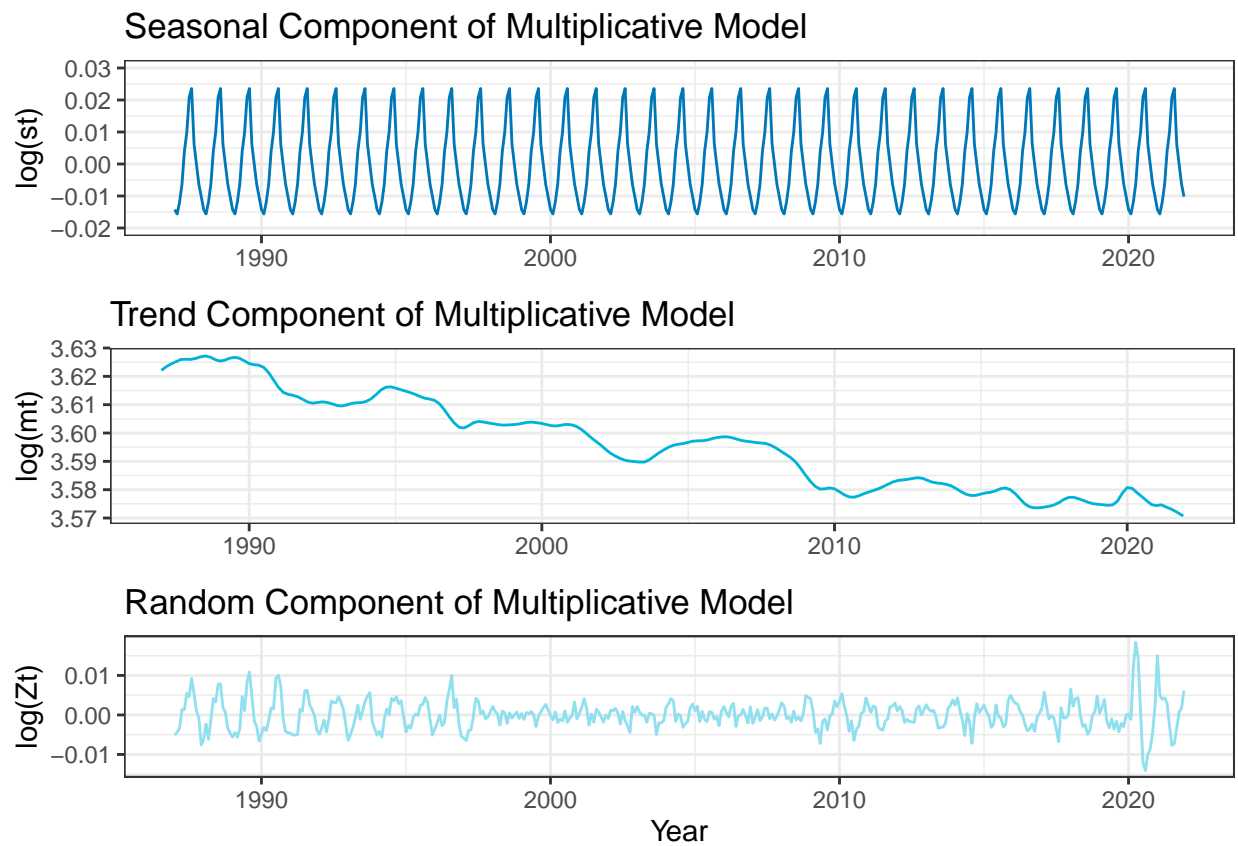
) +
ylim(-0.02, 0.03)

mp2 <- ggplot(data = fortify.zoo(trend), aes(x = Index, y = trend)) +
  geom_line(color = "#00b4d8", linewidth = 0.5) +
  theme_bw() +
  labs(
    title = "Trend Component of Multiplicative Model",
    y = "log(mt)",
    x = NULL
  )

zp2 <- ggplot(data = fortify.zoo(noise), aes(x = Index, y = noise)) +
  geom_line(color = "#90e0ef", linewidth = 0.5) +
  theme_bw() +
  labs(
    title = "Random Component of Multiplicative Model",
    y = "log(Zt)",
    x = "Year"
  )

grid.arrange(sp2, mp2, zp2, nrow = 3)

```



c) Fit a linear model to the trend component (you can use R function `lm()`).

- Write down the fitted model for the trend component.

Let $\hat{\beta}_0, \hat{\beta}_1 \in \mathbb{R}$ be the real-valued coefficient estimates of the true intercept and slope terms $\beta_0, \beta_1 \in \mathbb{R}$ respectively. Recalling that we are taking $\log(m_t)$, the fitted model is of the form.

$$\log(\hat{m}_t) = \hat{\beta}_0 + \hat{\beta}_1 t$$

The coefficients for (Intercept) ($\hat{\beta}_0$) and Time ($\hat{\beta}_1$) are found below:

```
# cast ts to data frame and rename
trend_df <- fortify.zoo(trend)
colnames(trend_df) = c("Time", "log_hours")
# fit model
trend_mod <- lm(log_hours~Time, data = trend_df)
# report coefficients
data.frame(value = trend_mod$coefficients)
```

```
##              value
## (Intercept)  6.744901235
## Time        -0.001570879
```

- Does the linear model provide evidence of a trend at the 95% confidence level?

To see if the linear model provides evidence of a trend component at the 95% confidence level, we would test the following pair of hypotheses at $\alpha = 0.05$:

$$H_0 : \hat{\beta}_1 = 0 \quad \text{against} \quad H_A : \hat{\beta}_1 \neq 0$$

In this instance, we would use the following test statistic:

$$T_{\text{obs}} = \frac{\hat{\beta}_1}{\text{se}(\hat{\beta}_1)} \sim t_{n-(k+1)}$$

Which takes a Student's t -distribution on $n - (k + 1) = 418$ degrees of freedom.

```
trend_summary = summary(trend_mod)

# get the standard error and estimated coefficients
se = trend_summary$coefficients[2, "Std. Error"]
est = trend_summary$coefficients[2, "Estimate"]

# calculate observed t where hypothesized value is zero
tobs = (est - 0)/se

# report
print(paste("Our observed test statistic is approximately",
            round(tobs, 2)))
```

```
## [1] "Our observed test statistic is approximately -77.81"
```

Then, we would calculate the p -value for the two-tailed test by finding:

$$p = 2 \min\{P(t_{n-(k+1)} > T_{\text{obs}}), P(t_{n-(k+1)} < T_{\text{obs}})\}$$

Which, in this case, is found as follows:

```
n = nrow(trend_df); k = 1

2 * min( pt( tobs, df = n - (k + 1), lower.tail = TRUE ),
         pt( tobs, df = n - (k + 1), lower.tail = TRUE ))

## [1] 8.217246e-251
```

Since our observed p -value of approximately $8.217 \times 10^{-251} \approx 0$ is less than our declared $\alpha = 0.05$, we would reject H_0 at the 95% level. There is statistically significant evidence to suggest that the $\hat{\beta}_1$ coefficient, the slope of the $\log(\hat{m}_t)$ model, is nonzero. Therefore, we conclude that the linear model provides evidence of a trend at the 95% confidence level.

- Without doing any further analysis, would you use this trend component to make predictions? Justify your answer using the linear model results and the trend component plot.

From the trend component plot, it appears that the trend m_t is non-linear as a function of time. Therefore, I would anticipate a pattern in the residual plot for the linear model $\log(\hat{m}_t)$. Despite the fact that the linear model was significant (see: previous question), it is very likely that there exists a better model (polynomial, etc.) to describe the relationship between $\log(\hat{m}_t)$ and t . If I were to make predictions, I would first try some other OLS models of different styles and, using cross-validation or otherwise, pick a method better than the “purely linear” approach taken here. However, despite the fact that it may not be the case that a linear model best suits $\log(\hat{m}_t)$, for the sake of prediction it is better to use this trend component rather than nothing at all. In short, it’s not the best model for $\log(\hat{m}_t)$, but it is better than nothing. I would prefer this trend component for predictions rather than ignoring trend altogether; however, it is probable that further analysis would find that a much better model exists for the trend component than this one.

d) Predict the monthly average values of the usual hours worked in Canada for the period from January 2022 to December 2023 using your seasonal decomposition model.

- Plot your predictions along with the actual observed values (on the same plot). Make sure to include a legend for your plot.

We’ll use the linear model from the previous question.

```
# extract model coefficients
beta_0 = trend_summary$coefficients[1, "Estimate"]
beta_1 = trend_summary$coefficients[2, "Estimate"]
# unique terms in the seasonal component give one period
period = unique(seasonal)
# use our linear model with time variable
mt_pred = beta_1*time(test) + beta_0
# we're predicting two periods into the future
st_pred = (rep(period, 2))
# then the predicted value is undoing the log transform
preds = exp(mt_pred + st_pred)
p3df <- data.frame(
```

```

    Time = as.numeric(time(test)),
    Actual = as.numeric(test),
    Predicted = preds)
# create custom-spaced month/date strings
date_strings <- unlist(lapply(2022:2023, function(year) {
  sapply(1:12, function(month) {
    if (month %% 6 == 0) paste(month, "/", year, sep = "")
    else ""
  })
}))
# specify additional parameters
date_strings[1] = "1/2022"
low_ylim = 34.5; up_ylim = 36.5

```

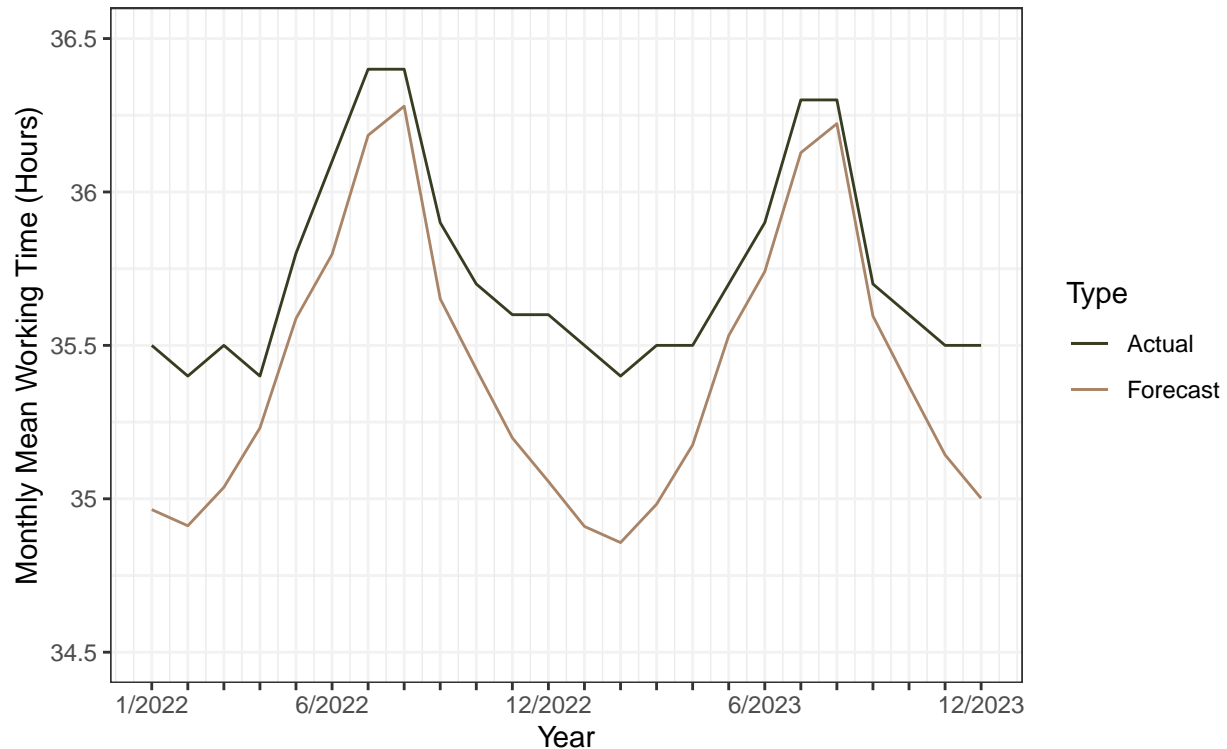
Now, we create a plot of the forecast vs. actual values.

```

p3 <- ggplot(p3df) +
  geom_line(aes(x = Time, y = Actual, color = "Actual")) +
  geom_line(aes(x = Time, y = Predicted, color = "Forecast")) +
  scale_color_manual(values = c("Actual" = "#373d20", "Forecast" = "#a98467"),
    name = "Type", labels = c("Actual", "Forecast")) +
  labs(
    title = "Monthly Average of Usual Hours Worked in Canada, with Forecast",
    subtitle = "Across all Industries from January 2022 to December 2023",
    y = "Monthly Mean Working Time (Hours)",
    x = "Year"
  ) + theme_bw() +
  scale_x_continuous(breaks = p3df$Time, labels = date_strings) +
  scale_y_continuous(limits = c(low_ylim, up_ylim),
    breaks = seq(low_ylim, up_ylim, by = 0.5),
    labels = seq(low_ylim, up_ylim, by = 0.5)) +
  theme(panel.grid.major = element_line(
    color = "grey95",
    linetype = "solid",
    linewidth = 0.5
  ),
  panel.grid.minor.y = element_line(
    color = "grey95",
    linetype = "solid",
    linewidth = 0.5
  ))
print(p3)

```


Monthly Average of Usual Hours Worked in Canada, with Forecast Across all Industries from January 2022 to December 2023




- Comment on the performance of your prediction method, explaining why or why not the method worked well for this data.

The prediction method seemed to perform decently well; specifically, in the capturing of the seasonal component of the time series. We see this fairly accurate performance of predicted seasons by the roughly-matching peak locations between the predicted values and the actual values in the test set.

However, there is certainly room for improvement. Consider the residual plot below:

```
p4df <- data.frame(
  Time = as.numeric(time(test)),
  Residual = as.numeric(test - preds)
)
p4 <- ggplot(p4df) +
  geom_line(aes(x = Time, y = Residual), color = "#780000") +
  labs(
    title = "Residual Plot of Actual Values against Forecast",
    subtitle = "Hours Worked Dataset (January 2022 to December 2023)",
    y = "Observed Residual",
    x = "Year"
  ) + theme_bw() + ylim(-0.05, 0.6) +
  geom_hline(yintercept = 0, lty = "dashed") +
  scale_x_continuous(breaks = p3df$Time, labels = date_strings)
print(p4)
```

Residual Plot of Actual Values against Forecast
Hours Worked Dataset (January 2022 to December 2023)



The residual plot shows a clear pattern, and the fact that the residuals are wholly positive implies that our model is consistently under-estimating the truth. The pattern to the residual plot is quite interesting, and may suggest that part of the reason the predictions are off is that there is some additional pattern in the data that we are not capturing with our current model.

It's difficult to say exactly what this "additional component" could be; however, there could be some secondary seasonal component (i.e. multiple-seasonality) that we are not detecting, or, as was discussed earlier, perhaps a purely linear model does not suit the trend component well.

- How could the prediction method be improved?

It could be improved in many ways. The first thing that comes to mind is the addition of error bounds, or any measure of uncertainty whatsoever. Having point estimate predictions isn't the greatest idea. Further, I think the precision of the method could be improved by attempting to implement non-linear functions of $\log(\hat{m}_t)$ and trying to use cross-validation methods (such as Lasso or GAMs) to better fit the trend component of the decomposition, ideally to improve both in-sample and out-of-sample performance. Finally, due to the pattern in the residual plot, I think there is some reason to suggest we should investigate more complex multi-seasonal models. It is plausible that there is some significant additional seasonality that we are not currently capturing with our decomposition that would improve the prediction method.

- As a statistician, what other information would you like to add to your forecasts in addition to the point forecasts you produced above?

As was mentioned in the previous question, it would be *extremely* wise to add confidence bounds / margins of error to our point forecasts. These would help quantify the uncertainty in our predictions and provide a much stronger model overall. Certain additional error quantification methods could also be helpful, such as

mean squared prediction error. But certainly without a doubt we need to quantify the uncertainty in our predictions somehow.

Task 2: Analyzing New York Temperature Data

Question (a)

Part 1: First, we read the data into R and create an R object called `dat`.

```
dat = read.csv("NY_Temperature_Data.csv")
```

Part 2: Now, we create a zoo object for daily Max Temperature.

NOTE: Our object is called `x`, but the input to `zoo()` is the daily maxima column from `dat`.

```
x = zoo(dat$TMAX, zoo::as.Date(dat$Date))
```

Part 3: Now, we make the monthly maxima time series, in an object called `monthly_max`.

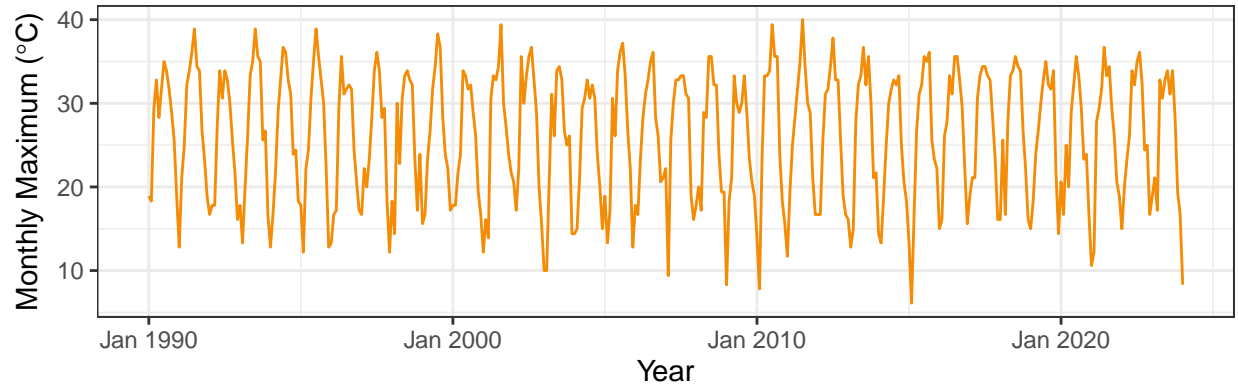
```
monthly_max = aggregate(x, as.yearmon, FUN=max)
```

Part 4: Now, we plot the monthly maximum temperature series. Since we're using `ggplot2` to create the plot, we'll use `fortify.zoo` that works alongside the `ggfortify` package to convert the `monthly_max` object to be `ggplot2` compliant.

```
p2df <- fortify.zoo(monthly_max)
p21 <- ggplot(p2df, aes(x = Index, y = monthly_max)) +
  geom_line(color = "#f48c06") + theme_bw() +
  labs(
    x = "Year",
    y = expression(paste("Monthly Maximum (", degree, "C)")),
    title = expression(paste(
      "Monthly Maximum Temperature (", degree, "C) in New York"
    )),
    subtitle = "Measured From 1990-2024, Sourced from NOAA"
  ) +
  coord_fixed(ratio = 0.275)
print(p21)
```

Monthly Maximum Temperature (°C) in New York

Measured From 1990–2024, Sourced from NOAA



Part 5: Comments on the observed features.

There are a few features we can observe. Firstly, the time series appears to be very high-frequency; that is, the period p is relatively short with respect to the overall time scale. However, there doesn't seem to be a change in amplitude over time, which, as discussed in the previous question, indicates the potential of using an additive model ($X_t = s_t + m_t + Z_t$) to describe this process. Finally, there doesn't immediately seem to be a high-magnitude trend component to this model; importantly, this doesn't imply that one doesn't exist. For example, we can see that the “troughs” (or lowest points) in the monthly maximum temperature seems to be decreasing slightly over time, implying that temperatures may be getting colder. There isn't conclusive evidence of a trend component simply from an initial observation; however, seeing if m_t is significant in this model certainly warrants investigation.

Question (c)

Fit a suitable seasonal decomposition model to the monthly data using the moving average smoothing (R function `decompose`) and plot the estimates of the trend, seasonal, and error components.

We start by using `zooreg` to convert to a time series.

```
month_ts <- ts(zooreg(monthly_max),  
              start = c(1990, 1),  
              end = c(2024, 1),  
              frequency = 12)
```

From our preliminary analytics in the previous part, we will decompose this time series as an additive model, i.e. $X_t = m_t + s_t + Z_t$, where $Z_t \sim \text{WN}(0, \sigma^2)$ for $t \in \mathbb{Z}$.

```

month_decomp <- decompose(month_ts)
mseasonal = month_decomp$seasonal
mtrend = month_decomp$trend
mnoise = month_decomp$random

```

For organization and ease of visualization, we de-contextualized the plot titles, simply mentioning it is for the Additive Model. However, it should be clarified that each plot title could be further specified to be “Component of Additive Model for Monthly Max Temperature in New York (1990-2024).”

```

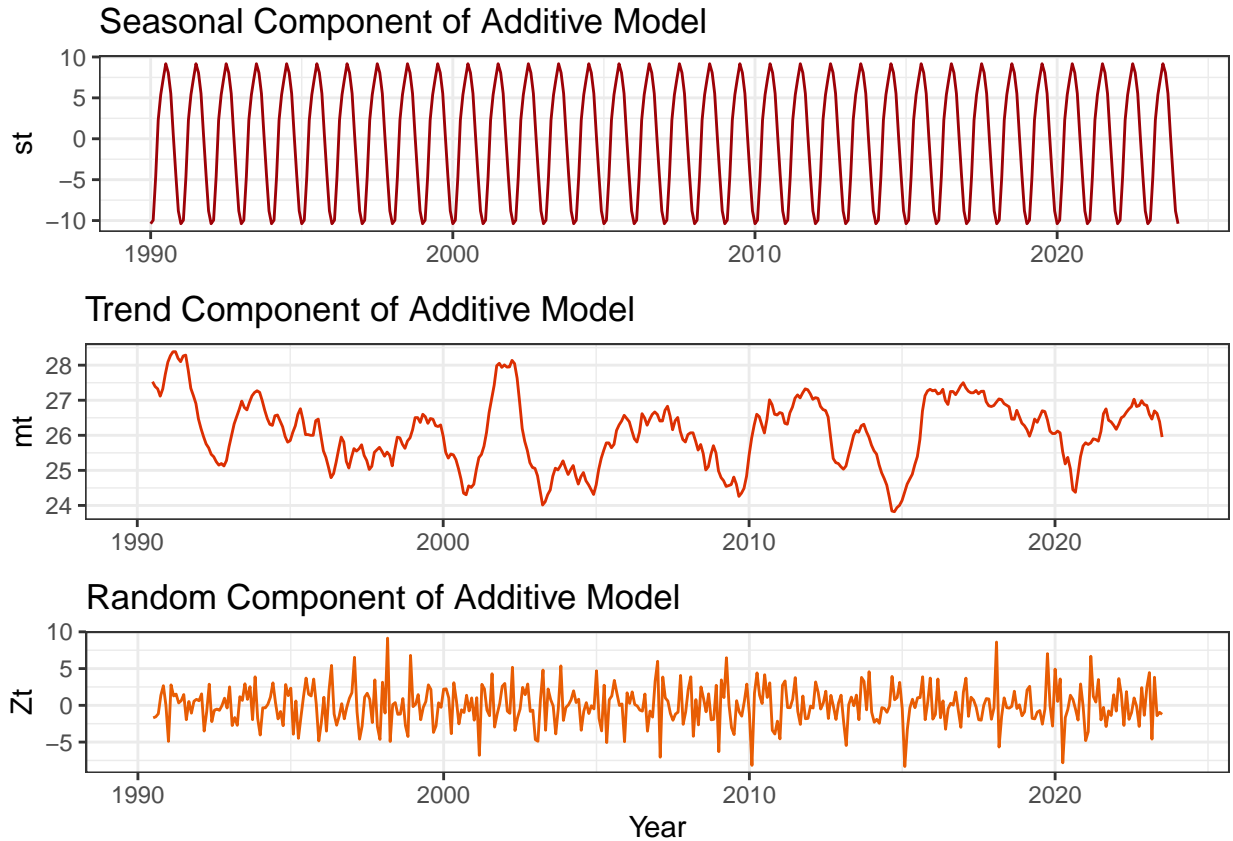
sp2c <- ggplot(data = fortify.zoo(mseasonal, na.fill = 0), aes(x = Index, y = mseasonal)) +
  geom_line(color = "#9d0208", linewidth = 0.5) +
  theme_bw() +
  labs(
    title = "Seasonal Component of Additive Model",
    y = "st",
    x = NULL
  ) + xlim(1990, 2024)

mp2c <- ggplot(data = fortify.zoo(mtrend, na.fill = 0), aes(x = Index, y = mtrend)) +
  geom_line(color = "#dc2f02", linewidth = 0.5) +
  theme_bw() +
  labs(
    title = "Trend Component of Additive Model",
    y = "mt",
    x = NULL
  ) + xlim(1990, 2024)

zp2c <- ggplot(data = fortify.zoo(mnoise, na.fill = 0), aes(x = Index, y = mnoise)) +
  geom_line(color = "#e85d04", linewidth = 0.5) +
  theme_bw() +
  labs(
    title = "Random Component of Additive Model",
    y = "Zt",
    x = "Year"
  ) + xlim(1990, 2024)

grid.arrange(sp2c, mp2c, zp2c)

```



Question (d)

Plot the correlogram for the deseasonalized series of monthly temperature maxima using the seasonal decomposition model you fit in part (c). Comment on the serial dependence of this series.

To de-seasonalize the data, we have to subtract the s_t component found in the previous section from the overall time series X_t found via aggregation. In terms of our R objects, the seasonal component is `mseasonal` and the time series observed (as a `ts` object) is `month_ts`. We compute their difference below:

```
deseasonalized = month_ts - mseasonal
```

Then, we plot the ACF $= \rho(h)$, using `na.pass` to compensate for the NA components of the `mtrend` object.

To informatively plot the sample acf, we need to determine how many time steps exist in our monthly data set so that we can accurately set the `lag.max` parameter. To do this, we let $\Delta t = t_2 - t_1$ be the time step between entries. For our data, $p = 12$, thus, $\Delta t = 1/12$ as calculated below.

```
# calculating what's one year in the TS
delta = time(deseasonalized)[2] - time(deseasonalized)[1]
```

Then, to determine the maximum lag value, we divide the total elapsed number of years in the time series by the time step size to find h_{\max} .

$$h_{\max} = \frac{(t_{\max} - t_{\min})}{\Delta t} = \frac{(2024 - 1990)}{(1/12)} = 408$$

```
# calculating what's the entire TS size (for max lag)
hmax = (2024-1990)/delta
```


Then, we can create the entire correlogram for our data.

```
d_acf <- acf(deseasonalized,
             na.action = na.pass,
             lag.max = hmax, plot = FALSE)
```

We know that if the sample acf values tend to be within $\pm 2/\sqrt{n}$, this is indicative of no serial dependence (i.e. simply *iid* noise.) These values are plotted below as horizontal dotted lines alongside the sample acf.

```
# create plotting dataframe
p2ddata = data.frame(
  h = 0:hmax,
  rh = d_acf$acf
)
# determine n
n = length(deseasonalized)
# determine significance bars for iid noise serial dependence
bars = c(-2, 2)/sqrt(n)
# create plot
p2d <- ggplot(p2ddata, aes(x = h, y = rh)) +
  geom_segment(aes(xend = h, yend = 0),
              color = "#606c38",
              linewidth = 0.5) +
  geom_hline(yintercept = bars[1], linetype = "dashed", col = "#283618")+
  geom_hline(yintercept = bars[2], linetype = "dashed", col = "#283618")+
  ylim(-0.2, 1)+
  geom_hline(yintercept = 0,
            linetype = "dashed",
            color = "darkgray") +
  labs(x = "Lag", y = "Autocorrelation",
       title = "Correlogram of De-Seasonalized Monthly Max Temperature Data",
       subtitle = "In New York 1990-2024, Data Sourced from NOAA") +
  theme_bw()
print(p2d)
```

Correlogram of De-Seasonalized Monthly Max Temperature Data In New York 1990–2024, Data Sourced from NOAA



Comments

Visually, it appears that the de-seasonalized data seem to be predominantly white noise (i.e. no serial dependence) that is decaying very slowly over time (as we'd expect.) Further, the majority of the observations fall below the $\pm 2/\sqrt{n}$ bars.

To verify this observation, we'll let Υ be the proportion of $\rho(h)$ values greater than the $\pm 2/\sqrt{n}$ threshold.

We'll compute the proportion as follows:

$$\Upsilon = \frac{1}{(h_{\max} + 1)} \sum_{i=0}^{h_{\max}} \mathbb{1}\left(|\rho(h_i)| \geq \frac{2}{\sqrt{n}}\right)$$

In our R code, we compute the value below:

```
# same as mean( abs(d_acf$acf) > bars[2] )
p_greater = sum( abs(d_acf$acf) > bars[2] ) / (hmax + 1)
# show as percent
(p_greater * 100)
```

```
## [1] 0.9779951
```

It appears that only approximately 0.98% the sample acf values lie above the random noise threshold (noting that this include $\rho(0) = 1$ that is always above this threshold.) It seems overwhelmingly likely that there is no serial dependence to the de-seasonalized data.

In addition to the computation from the threshold above, we don't see any periodicity to the correlogram that would indicate any "left over" seasonality, nor do we see a slow decay of high-magnitude autocorrelation

values on the same side of the mean that would indicate a trend. Hence, we can conclude that it is likely that there is no serial dependence present in the de-seasonalized data.

Task 3: Conducting a Simulation Study on the Autocorrelation Coefficient

Recall that for a time series of length n , from a white noise process, for n sufficiently large the sample autocorrelation coefficient at lag h approximately follows a normal distribution with mean $-1/n$ and variance $1/n$.

Part (i)

Simulate a time series of length $n = 2000$ from a white noise process with $Z_t \sim N(0, 1)$ (function `rnorm()`).

Below, we compute $n = 2000$ realizations of a $N(0, 1)$ random variable. We will let $\mathcal{Z}_1 = \{z_{1,1}, z_{1,2}, \dots, z_{1,n}\}$ be this set of realizations where $z_{i,j}$ is the j -th realization from the i -th vector of simulated values, for $i \in [1, m]$ and $j \in [1, n]$.

```
z_1 = rnorm(n = 2000, mean = 0, sd = 1)
```

Part (ii)

Evaluate the sample autocorrelation coefficient r_h at lag h for $h = 1$ and $h = 2$. Store these values.

We note that in this case we slice the vector of observed $\rho(h)$ values at index $h + 1$, since R language vectors begin at index 1. Hence, `$acf[1]` would give $\rho(0) = 1$ which isn't what we want here.

```
r_1 = acf(z_1, plot = FALSE)$acf[2]
r_2 = acf(z_1, plot = FALSE)$acf[3]
round(c(r_1, r_2), 4)
```

```
## [1] -0.0223 -0.0060
```

Part (iii)

Repeat steps (i) and (ii) $m = 8000$ times;

Generate 8000 time series of length n and for each of them compute r_1 and r_2 . You should now have two vectors of length m with estimates r_1 and r_2 .

The code below uses the following algorithm:

Algorithm

Repeat the following $\forall i \in [1, m]$ (assuming *iid* white noise Z_t):

- 1.) Let $\mathcal{Z}_i = \{z_{i,1}, z_{i,2}, \dots, z_{i,n}\}$ be n realizations from an $N(0, 1)$ random variable.
- 2.) Let $\rho_i(h)$ be the sample acf from \mathcal{Z}_i .
- 3.) Let $r_{i,1} = \rho_i(1)$ be the auto-correlation of lag $h = 1$ from this iteration.
- 4.) Let $r_{i,2} = \rho_i(2)$ be the auto-correlation of lag $h = 2$ from this iteration.
- 5.) Append $r_{i,1}$ to vector \mathcal{R}_1 and append $r_{i,2}$ to vector \mathcal{R}_2 .

Result: After m iterations, we will have $\mathcal{R}_1 = \{r_{1,1}, r_{2,1}, \dots, r_{m,1}\}$ and $\mathcal{R}_2 = \{r_{1,2}, r_{2,2}, \dots, r_{m,2}\}$ such that $|\mathcal{R}_1| = |\mathcal{R}_2| = m$, as required.

This algorithm is implemented directly in the code below:

```
set.seed(11292023)
m = 8000; n = 2000
simulation = sapply(1:m,
  # Algorithm
  function(i){
    z_i = rnorm(n)
    rho_i = acf(z_i, plot = FALSE)
    r_i1 = rho_i$acf[2]
    r_i2 = rho_i$acf[3]
    c(r_i1, r_i2)
  })
# Results
results = data.frame(t(simulation))
colnames(results) = c("R1", "R2")
head(results)
```

```
##           R1           R2
## 1 -0.019810489  0.0177181321
## 2  0.009318669  0.0221394680
## 3  0.022568568 -0.0273215956
## 4 -0.015705478  0.0059205404
## 5  0.009394935 -0.0008523259
## 6  0.019431549  0.0339157580
```

Part (iv)

Now, we summarize the results of the simulation study in the following three sections.

Part (iv.1.) We will let $\mathbb{E}(\mathcal{R}_1)$ and $\mathbb{E}(\mathcal{R}_2)$ be the mean of the simulated r_1 and r_2 values from the study, respectively. Similarly, we let $\text{var}(\mathcal{R}_1)$ and $\text{var}(\mathcal{R}_2)$ be the variance of the r_1 and r_2 values from the study, respectively.

These values are computed and summarized in the table below.

```
tableDF <- data.frame(names = c(colnames(results), "Theorized Value"),
  means = c(sapply(results, function(r) {sprintf("%.2e", mean(r))}),
    -1/n),
  vars = c(sapply(results, function(r) {sprintf("%.2e", var(r))}),
    1/n)
)
colnames(tableDF) = c("Lag", "Mean", "Variance")

gt_table <- gt(tableDF) %>%
  tab_header(title = "Simulation Results",
    subtitle = md("From `m = 8000` Iterations")) %>%
  tab_style(
    style = cell_text(weight = "bold"),
    locations = cells_column_labels(columns = everything())
  ) %>% tab_options(
    stub.border.width = px(2),
```

```

    stub.border.color = "black"
  ) %>%
  gt_add_divider(columns = "Lag", style = "solid", col = 'grey70')

(gt_table)

```

Simulation Results
From $m = 8000$ Iterations

Lag	Mean	Variance
R1	-6.78e-04	4.99e-04
R2	-4.84e-04	5.15e-04
Theorized Value	-5e-04	5e-04

Part (iv.2.) In two separate figures plot the two histograms for the sample of r_1 and r_2 values from the simulation study (function `hist()`) add the smoothed version of the histogram (function `density()`) and the theoretical asymptotic normal density (function `dnorm()`). Make sure your plots are well-presented including a suitable title, axes labels, curves of different type or colour, and a legend.

We'll begin with the plot of observed r_1 values.


```

x <- seq(from = -0.05, to = 0.05, length.out = m)
p3b1 = ggplot(results, aes(x = R1)) + theme_bw() +
  geom_histogram(
    aes(y = after_stat(density)),
    binwidth = 0.0065,
    colour = "#555B6E",
    fill = "#fdfdff",
    position = "identity"
  ) +
  geom_density(aes(color = "Observed Density")) +
  labs(x = expression(paste("Simulated Values of ", r[1] )),
    y = "Density",
    color='Legend:',
    title = expression(paste("Histogram of Simulated ", r[1], " Values")),
    subtitle = "With Density Curve and Theoretical Asymptotic Normal Curve") +
  geom_line(aes(
    x = x,
    y = dnorm(x, mean = -1 / n, sd = sqrt(1 / n)),
    color = "Theorized"
  )) +
  scale_color_manual(values = c(
    "Observed Density" = "#594436",
    "Theorized" = "#89B0AE"
  )) +
  theme(legend.position = "top",
    panel.grid.minor = element_line(
      color = "grey90",
      linetype = "dashed",
      linewidth = 0.5
    ))
print(p3b1)

```

Histogram of Simulated r_1 Values

With Density Curve and Theoretical Asymptotic Normal Curve



Now, we'll plot the simulated r_2 values.

```
p3b2 = ggplot(results, aes(x = R2)) + theme_bw() +
  geom_histogram(
    aes(y = after_stat(density)),
    binwidth = 0.0065,
    colour = "#555B6E",
    fill = "#ffdfd",
    position = "identity"
  ) +
  geom_density(aes(color = "Observed Density")) +
  labs(x = expression(paste("Values of ", r[2] )),
    y = "Density",
    color='Legend:',
    title = expression(paste("Histogram of Simulated ", r[2], " Values")),
    subtitle = "With Density Curve and Theoretical Asymptotic Normal Curve") +
  geom_line(aes(
    x = x,
    y = dnorm(x, mean = -1 / n, sd = sqrt(1 / n)),
    color = "Theorized"
  )) +
  scale_color_manual(values = c(
    "Observed Density" = "#432534",
    "Theorized" = "#c44900"
  )) +
  theme(legend.position = "top",
```

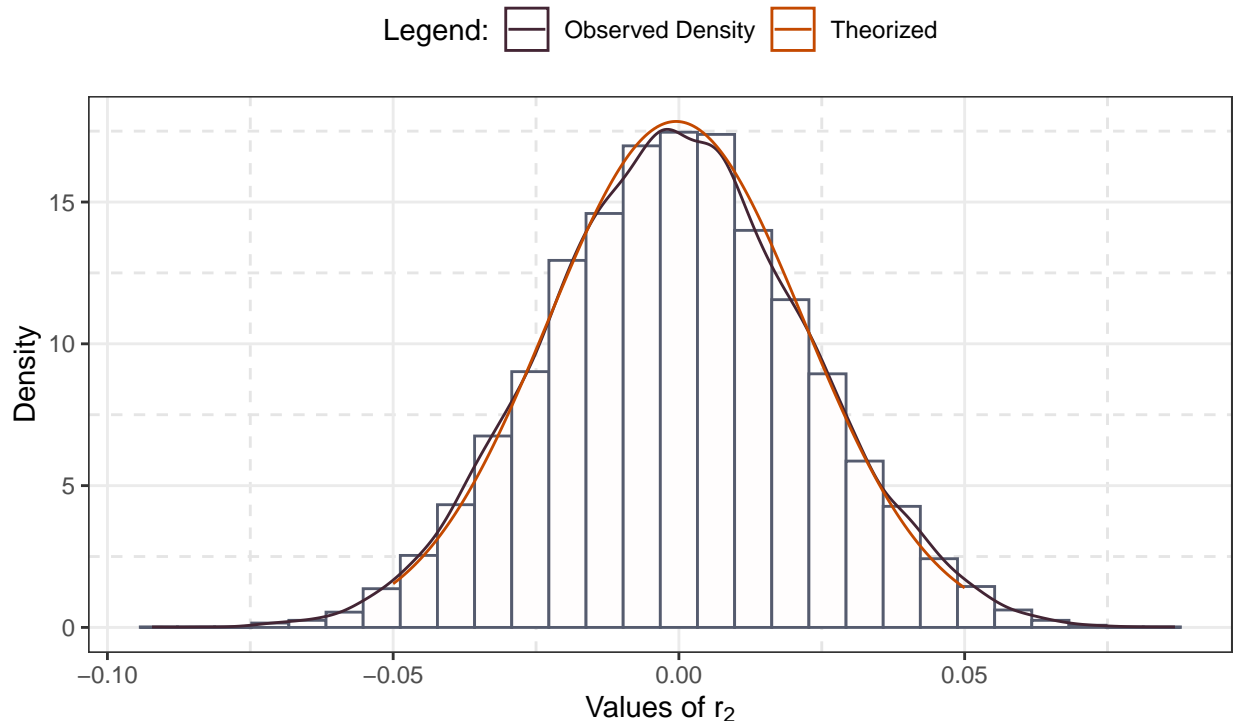
```

panel.grid.minor = element_line(
  color = "grey90",
  linetype = "dashed",
  linewidth = 0.5
))
print(p3b2)

```

Histogram of Simulated r_2 Values

With Density Curve and Theoretical Asymptotic Normal Curve



Part (iv.3.)

- Comment whether there is an agreement between the empirical estimates of the bias, variance, and sampling density of the estimator of the autocorrelation at lag h and their theoretical approximation.

Visually, it appears that there is an agreement between the sampling density of the estimator of the autocorrelation with their theoretical normal values at both $h = 1$ and $h = 2$.

Further, it seems as if the observed density curve and theorized normal densities are centered at approximately the same location for both lag values.

Finally, we can see that the two curves follow approximately the same dispersion/spread as the theoretical normal density. It appears slightly heavier-tailed than the theorized distribution; however, not to an extreme degree.

In short, from our empirical observations based on the plots in the previous question, it seems likely that both the true r_1 and r_2 values follow the theoretical approximation.

However, we may wish to also analytically verify these claims. To do so, we will construct confidence intervals for the mean and variances at $\alpha = 0.05$, i.e. 95%. We will begin by constructing pairwise confidence intervals

for the means of r_1 and r_2 to determine if $-1/n$ is in these intervals. To adjust for family-wise error rate, we'll apply the Bonferroni correction and use $\alpha' = 0.05/2$, in other words constructing 97.5% intervals.

We'll construct the CI s for r_1 and r_2 by the following methodology, using sample size $N = m$ and a t -Distribution with $m - 1$ degrees of freedom (noting for values of N this large, the t distribution is approximately standard normal):

$$\text{For } h \in [1, 2], \quad CI_h(\alpha) = \left(\mathbb{E}(\mathcal{R}_h) \pm t_{\alpha/2; (m-1)} \frac{\text{se}(\mathcal{R}_h)}{\sqrt{m}} \right)$$

The code to compute the two simultaneous intervals is below

```
alpha = 0.05
alpha_prime = alpha/2
cis = data.frame( apply(results, MARGIN = 2,
  function(r){
    rbar = mean(r)
    se_i = sd(r) / sqrt(m)
    moe = qt(alpha_prime/2, df = m-1)*se_i
    return(rbar + c(moe, -moe))
  })
rownames(cis) = c("Lower", "Upper")
signif(cis, 3)
```

```
##           R1           R2
## Lower -0.001240 -1.05e-03
## Upper -0.000118  8.49e-05
```

As the intervals show, the value of $-1/n = -5 \times 10^{-4}$ lies within both the r_1 and r_2 intervals. Hence, we know that the corresponding hypotheses for equality of means would fail to reject H_0 at the 95% level. In other words, there is statistically significant evidence (at $\alpha = 0.05$) to suggest that the means of the simulated values follow the theoretical approximations, for this particular seed. Further, from the results of this hypothesis test, we can conclude that there is not statistically significant evidence at $\alpha = 0.05$ to suggest that either lag $h = 1$ or $h = 2$ have bias relative to their theoretical distributions.

$$\forall h \in [1, 2], \quad \text{bias}(r_h) = \mathbb{E}(\mathcal{R}_h) - \left(-\frac{1}{n} \right) = 0$$

We can follow a similar methodology to determine whether or not our simulations follow the theorized variance of $\sigma^2 = 1/n$. Again, we will construct two simultaneous confidence intervals adjusting for FWER, noting that for variance the intervals are constructed as follows (where m , α , \mathcal{R} and h are as previously declared):

$$\text{For } h \in [1, 2], \quad CI_h(\alpha) = \left(\frac{\text{var}(\mathcal{R}_h)(m-1)}{\chi^2_{1-(\alpha/2); (m-1)}}, \frac{\text{var}(\mathcal{R}_h)(m-1)}{\chi^2_{\alpha/2; (m-1)}} \right)$$

```
civs = data.frame( apply(results, MARGIN = 2,
  function(r){
    numer = var(r)*(m-1)
    ci_v = c( numer / qchisq(1-alpha_prime/2, df = (m - 1)),
              numer / qchisq(alpha_prime/2, df = (m - 1)))
    return(ci_v)
  })
rownames(civs) = c("Lower", "Upper")
signif(civs, 3)
```

##	R1	R2
## Lower	0.000482	0.000497
## Upper	0.000517	0.000534

As the intervals show, the value of $1/n = 5 \times 10^{-4}$ lies within both the r_1 and r_2 variance confidence intervals. Hence, we know that the corresponding hypotheses for equality of variances would fail to reject H_0 at the 95% level. In other words, there is statistically significant evidence (at $\alpha = 0.05$) to suggest that the variance of the simulated values follow the theoretical approximations, for this particular seed.

Putting all this together, the consistency of the confidence intervals with the theoretical approximations give us some important analytical credence to our initial empirical observations. We can conclude that it is overwhelmingly likely true that the simulated data follow the theorized distribution.