

Stat 443: Time Series and Forecasting

Assignment 3: Time Series Models

Caden Hewlett

March 20, 2024

Question 1: El Niño Forecasting

The file `NIN034.csv` contains the monthly El Niño 3.4 index from 1870 to 2023. The El Niño 3.4 index represents the average equatorial sea surface temperature (in degrees Celsius) from around the international dateline to the coast of South America.

Part a

Perform exploratory data analysis.

Part a.1

Import the data into R and create a time-series object for the El Niño 3.4 index.

```
# import data
# setwd('C:/Users/caden/OneDrive/Desktop/STAT_443/STAT443')
df <- read.csv("NIN034.csv")

# raw data is in a bad format, need to pivot it without years
dflong <- df %>%
  pivot_longer(cols = -Year, names_to = "Month", values_to = "Value")
# re-translate into dates with month abbreviations
dflong <- dflong %>%
  mutate(Date = make_date(Year, match(Month, month.abb), 1))
# then, format into a time series!
nino_ts <- ts(dflong$Value,
             start = c(1870, 1), frequency = 12)

# time series length = data frame length = 154 years * 12 months per year
stopifnot(all.equal(length(nino_ts), nrow(df)*12, 154*12))
```

Break the time series object into a training and test set. You can use the function `window()` on a `ts` object to split the data. Let the training set be from January 1870 to December 2021, and let the test set start in January 2022 and end in November 2023.

```
# split data into train and test
nino_train = window(nino_ts, start = c(1870, 1), end = c(2021, 12))
nino_test = window(nino_ts, start = c(2022, 1), end = c(2023, 11))
# verify split
stopifnot(all.equal(
  length(nino_test)+length(nino_train), length(na.remove(nino_ts))
))
```

```
)  
)
```

Part a.2

Plot the training data as well as its acf and pacf.


Training Data

```
p1data = fortify.zoo(nino_train)
```

```
print(p1)
```

El Nino 3.4 Index from Jan. 1870 to Dec. 2021

Index Represents Average Equatorial Sea Surface Temperature (in °C)




Autocorrelation Function

```
p2data = data.frame(  
  h = 0:36,  
  rh = acf(nino_train, plot = FALSE, lag.max = 36)$acf  
)
```

```
print(p2)
```

Correlogram of El Nino 3.4 Index

Using Training Data Spanning Jan. 1870 to Dec. 2021




Partial Autocorrelation Function

```
p3data = data.frame(  
  h = 1:36,  
  rhh = pacf(nino_train, plot = FALSE, lag.max = 36)$acf  
)
```

```
print(p3)
```

Partial ACF of El Nino 3.4 Index

Using Training Data Spanning Jan. 1870 to Dec. 2021




Comments

From the ACF, it is very likely that the original series exhibits a trend. This is due to the fact that there is a clear sinusoidal component of the ACF that is not dampening at a high rate (as we could potentially expect in models with AR components with $\alpha_i < 0$ for $i \in [1, p]$).

In fact, we can replicate the pattern in the sample ACF quite closely by simply using a sine function alongside random noise. Let $\{Z_t\}_{t=1}^{500} \stackrel{\text{iid}}{\sim} N(0, 1)$ and let $X_t = Z_t + \sin(t/2)$, for $t \in \{1, 2, \dots, 500\} \subset \mathbb{N}$. A plot of this artificial additive seasonal model is below:

```
set.seed(443); Z = rnorm(500); t = seq(1:500)
acf(Z + sin(t/2), main = "Artificial ACF", lag.max = 36)
```

Artificial ACF



As you can see, this artificial ACF closely matches our observed sample ACF from earlier. As such, it is very likely that there exists some seasonal term s_t in the time series of sea surface temperature data. We would intuitively anticipate a seasonal component to these data, as it is very likely that sea temperatures vary over the course of the year due to seasons as measurements are coming from the same area (the international dateline to the coast of South America.) Moreover, from a more informal observational perspective, a plot of the original series data doesn't seem to support a non-constant seasonal amplitude - as in there's no changing peak heights over time - which may indicate an additive seasonality rather than multiplicative.

In addition, neither the plot of the data nor the ACF seem to indicate a significant trend in the data. If there were a trend component m_t , we would see some consistent change over time in the overall direction of the series beyond simple seasonality. It doesn't seem like there is anything like that in this series; **however**, a purely visual analysis isn't comprehensive so this doesn't mean that a trend component doesn't exist.

Finally, to determine whether or not the series is stationary, we recall the definition of a weakly stationary stochastic process. Specifically, we will consider the first property of weak stationarity - that the mean is constant. We defined this formally in Assignment 1 previously, and was given similar to the following:

$$\text{Weak Stationarity Property One: } \exists \mu \in \mathbb{R} \text{ s.t. } \forall t \in \mathbb{Z}, \mathbb{E}(X_t) = \mu$$

The presence of *either* a seasonal component or a trend causes a contradiction, as the expected value of the series becomes some function of t (and hence cannot be a constant.)

Thus, for our particular series, we have the following implication, letting s_t be the seasonal component, $c \in \mathbb{R}$ be some real constant and $f(t)$ be a non-constant real function (as in, it changes with t .) Then,

$$\exists s_t \implies \forall t \mathbb{E}(X_t) = f(t) + c \implies \nexists \mu \in \mathbb{R} \text{ s.t. } \forall t \mathbb{E}(X_t) = \mu \therefore X_t \text{ is not stationary. } \square$$

In short, due to the presence of a trend, the time series for the training data is not stationary. It should be noted that not much can be concluded with results from the PACF, as the series is non-stationary.

Part b

Forecast sea surface temperature for 2022 and 2023 using the Box-Jenkins method and the data from 1870-2021.

Part b.1.


Remove any seasonal variation and trend from the training data, if there is any, using the `stl` function in R.

We define an object `nino_stl` using `loess` decomposition, and plot each of S, T and L (seasonal, trend and irregular) of the original series below:

```
# perform loess decomposition
nino_stl = stl(nino_train, s.window = "periodic")
# cast to data frame
decomposition = as.data.frame(nino_stl$time.series)
decomposition$time = as.numeric(time(nino_train))
```

Note that the code to plot the decomposition with `ggplot` became quite long, so I made the code cell hidden. The complete plots are provided below:

```
grid.arrange(pb1, pb2, pb3, ncol = 1)
```




Now, we plot the filtered data set, as well as its acf and pacf. It doesn't seem like there is a significant trend component, as our original observations also showed, so we merely de-seasonalize the data. The removal of s_t and the plot of the filtered data is below:

Filtered Data

```
# remove the season
nino_filtered = nino_train - decomposition$seasonal
# prepare the filtered series for plotting
pb4data = fortify.zoo(nino_filtered)

print(pb4)
```

Deseasonalized El Nino 3.4 Index from Jan. 1870 to Dec. 2021
Seasonal Component Identified via Loess Decomposition




Autocorrelation

```
# compute acf values for lags
p5bdata = data.frame(
  h = 0:45,
  rh = acf(nino_filtered, plot = FALSE, lag.max = 45)$acf
)

print(pb5)
```

Sample ACF of Deseasonalized El Nino 3.4 Index Using Training Data Spanning Jan. 1870 to Dec. 2021




Partial Autocorrelation

```
p6bdata = data.frame(  
  h = 1:45,  
  rhh = pacf(nino_filtered, plot = FALSE, lag.max = 45)$acf  
)
```

```
print(pb6)
```


Partial ACF of Deseasonalized El Nino 3.4 Index Using Training Data Spanning Jan. 1870 to Dec. 2021



Part b.2.

In the work to follow, we assume the de-seasonalized time series is stationary, such that the ACF and PACF can be used for model identification.

Now, using the standard graphical tools, let's select two models from the pure MA and pure AR family of models for the filtered data set.

Since the ACF follows a (slightly complicated) exponentially decaying sinusoidal pattern, it seems quite likely that the best candidate models will be from the AR family. If this were a pure MA process for some order q , we would see the ACF sharply cut off after the value of q is reached. Instead, the ACF slowly decays over time; this is indicative of an AR component to the series. Hence, we'll consider models from the AR family of models.

Now, given that we are attempting to find p under the presumption that the series follows some pure AR process, we can utilize the PACF to attempt to find a suitable p . Unfortunately, it isn't immediately evident which term p in the PACF satisfies the condition we commonly use that $\forall p > k, |\hat{\alpha}_{kk}| < 2/\sqrt{n}$.

We will consider two potential candidates, the first being AR(1) process. The justification for deciding upon $p = 1$ is based on the fact that $\hat{\alpha}_{11}$ is significantly larger than all other PACF values. While it is true that $\exists k > 1$ s. t. $|\hat{\alpha}_{kk}| \geq 2/\sqrt{n}$, there are no other values in the PACF greater than $\hat{\alpha}_{11}$. In other words, the PACF may not decline into noise; however, $\forall k \in [2, n] \subset \mathbb{Z}, \hat{\alpha}_{11} > \hat{\alpha}_{kk}$. So, while the PACF doesn't cut off into white noise after $p = 1$, it does cut off to significantly smaller values.

The second candidate is an AR(6) process. The reasoning is due to the fact that after $\hat{\alpha}_{66}$, the vast majority of observed $\hat{\alpha}_{kk}$ values for $k > 6$ seem to be below the white noise threshold. Critically, there isn't an easily distinguishable pattern to those $\hat{\alpha}_{kk}$ values greater than $\pm\sqrt{2}/n$ after PACF lag 6. This could mean that for $k > 6$ there is simply noise (as we want), however the data could have higher noise values than those

commonly assumed by the W.N. threshold. It's important to recall that $\pm\sqrt{2}/n$ is more of a “rule of thumb” rather than a strict rule.

With all of this mind, we fit the models below:

```
# ar(1) vs. ar(6)
ar1 = arima(nino_filtered, order = c(1, 0, 0))
ar6 = arima(nino_filtered, order = c(6, 0, 0))
```

Then, we extract the coefficients and show them in the table below, rounding each to three decimal places:

```
ar_fits = t(round(data.frame(
  AR1 = c(ar1$coef[-1], head(ar1$coef, -1), rep(NA, times = 5), ar1$sigma2),
  AR6 = c(tail(ar6$coef, 1), head(ar6$coef, -1), ar6$sigma2)), 3))
# report with nice latex column names
kable(ar_fits, "latex", escape = FALSE, col.names = var_col_names,
  caption = "Fitted AR Coefficients") %>%
kable_styling(latex_options = "hold_position") %>%
kable_styling(position = "center")
```

Table 1: Fitted AR Coefficients

	$\hat{\mu}$	$\hat{\alpha}_1$	$\hat{\alpha}_2$	$\hat{\alpha}_3$	$\hat{\alpha}_4$	$\hat{\alpha}_5$	$\hat{\alpha}_6$	$\hat{\sigma}^2$
AR1	26.955	0.926	NA	NA	NA	NA	NA	0.085
AR6	26.954	0.920	0.129	-0.074	0.009	-0.016	-0.087	0.080

Part b.3.

Compare the AIC values of your two models. Which would you pick based on AIC?

First, we extract the aAICc values of each fit.

```
kable(
  t(data.frame(
    AR1 = round(ar1$aic, 3),
    AR6 = round(ar6$aic, 3)),
  col.names = c("Candidate Model", "AIC"),
  caption = "Akaike Information Criterion"
) %>%
kable_styling(latex_options = "hold_position") %>%
kable_styling(position = "center")
```

Table 2: Akaike Information Criterion

Candidate Model	AIC
AR1	687.816
AR6	596.045

Let $AIC(m)$ be the AIC for models $m \in M$, where $M = \{AR(1), AR(6)\}$. Generally, models with smaller AIC values are preferred.

So, we can select the best model m_{best} with the following logic:

$$m_{\text{best}} = \underset{m \in M}{\operatorname{argmin}} \{AIC(m)\}, \text{ where } M = \{AR(1), AR(6)\}$$

Then, using the reported values above, we can select m_{best} as follows:

$$\begin{aligned}
m_{\text{best}} &= \operatorname{argmin}_{m \in M} \{ \text{AIC}(m) \} \\
m_{\text{best}} &= \operatorname{argmin}_{m \in M} \{ \text{AIC}(\text{AR}(1)), \text{AIC}(\text{AR}(6)) \} \\
m_{\text{best}} &= \operatorname{argmin}_{m \in M} \{ 687.816, 596.045 \} \\
m_{\text{best}} &= \text{AR}(6)
\end{aligned}$$

So, by AIC selection, we prefer the AR(6) model.

Part b.4.

We recall that the third plot of `tsdiag` are the p -values of the Ljung-Box Version of the Portmanteau Lack-of-Fit test, where the following hypotheses are being compared:

H_0 : The Model is Well-Specified. v.s. H_A : The Model is Not Well-Specified.

Where the p -value is found via the χ^2 distribution and the test statistic, given by the following for a pure AR model (i.e. $q = 0$)

$$Q_2 = n(n+2) \sum_{i=1}^h \frac{\hat{\rho}_z(h)^2}{n-h} \sim \chi_{h-p}^2$$

To get a better understanding of the plots to follow, we compute the Q_2 values for $h \in [1, 10] \subset \mathbb{Z}$ below.


```

n = length(nino_filtered)
lb_stats = matrix(0, ncol = 2, nrow = 10)
M = list(ar1, ar6)
for(i in 1:length(M)){
  # get i-th model
  m = M[[i]]
  # get rho(h) values of residuals
  rh_m = acf(m$residuals, lag.max = 10, plot = F)$acf
  # compute Ljung-Box stat
  lb_ar = sapply(2:11,
    function(h){
      Q = n*(n+2)*sum( rh_m[2:h]^2 / (n - h))
      p = pchisq(Q, df = h - 1, lower.tail = F) # p = 1 for AR(1)
      p
    })
  lb_stats[, i] = lb_ar
}

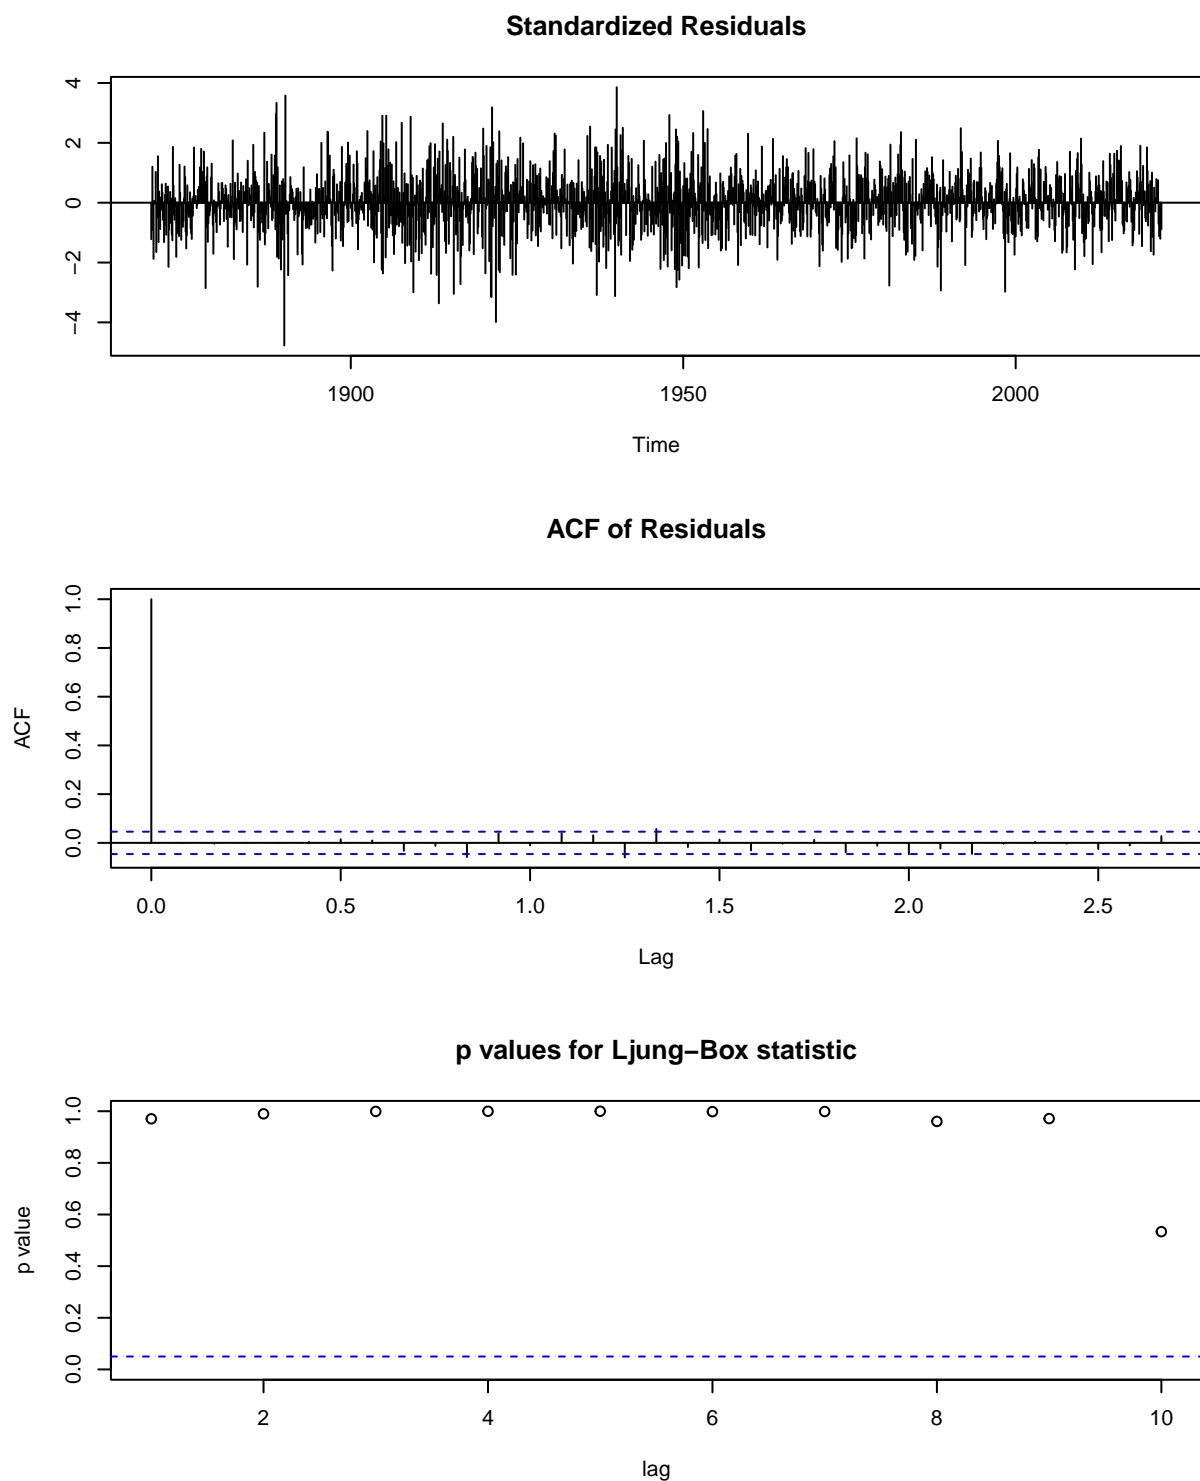
```

Now, we use the `tsdiag` function in R to plot diagnostics for your two models.

```
tsdiag(ar1)
```



```
tsdiag(ar6)
```



What do you observe?

We see that the AR(1) model has Ljung-Box p -values below the $\alpha = 0.05$ significance threshold $\forall h > 1$, which is confirmed by out `lb_stats` object, as you can see below:

```
which(lb_stats[, 1] < 0.05)
```

```
## [1] 2 3 4 5 6 7 8 9 10
```

This tells us that we reject H_0 in favour of the alternative that the model is ill-fitting for the AR(1) case in the majority of scenarios. This implies that it may be likely that the AR(1) model is not well-specified for these data. This theory is supported by the ACF of the residuals, which appears to have a vague sinusoidal pattern to it, as well as having values above $\pm 2/\sqrt{n}$ for quite a few of the lags. All in all, the `tsdiag` results for the AR(1) model shed doubt on our confidence of this candidate being well-fitting for these data.

To contrast, the Ljung-Box statistic values for the AR(6) model are consistently well above $\alpha = 0.05$, as shown in the table below (for the first five lags to conserve space):

```
kable(t(data.frame(
  Lag = 1:5,
  AR1 = format(lb_stats[1:5, 1], scientific = TRUE, digits = 5),
  AR6 = format(lb_stats[1:5, 2], scientific = TRUE, digits = 5))),
  caption = "Ljung-Box Portmanteau Test  $p$ -Values for Varying Lags") %>%
  kable_styling(latex_options = "hold_position") %>%
  kable_styling(position = "center")
```

Table 3: Ljung-Box Portmanteau Test p -Values for Varying Lags

Lag	1	2	3	4	5
AR1	1.4506e-01	2.6950e-10	4.0126e-11	2.1668e-12	1.8243e-12
AR6	9.7066e-01	9.8999e-01	9.9923e-01	9.9995e-01	9.9998e-01

This shows that the probability of observing the computed Q_2 values given H_0 that the model is well-specified is much higher for the AR(6) model than the AR(1) alternative. Further, the ACF of the residuals of the AR(6) model seem to quickly decay towards white noise, and shows far less extremity and patterning than the AR(1) ACF. So, if we were deciding the best candidate m_{best} purely from these results, we would again prefer the AR(6) model.

Part b.5.

Predict the sea surface temperature for Jan 2022 through Nov 2023 using both candidate models. Use this to calculate the mean squared prediction error (MSPE) for both models.

We again let $m \in M$ be the candidate models, for $M = \{\text{AR}(1), \text{AR}(6)\}$. We let $\{y_t\}_{t \in \mathbb{Z}}$ be the observed values from the test set. We let $\hat{y}_t^{(m)}(\ell)$ be the ℓ -step ahead Box-Jenkins forecast (for $\ell \in \mathbb{Z}$) of model m . Let N_ℓ be the total number of steps ahead. In our case, this is 23 for both.

The set of MSPE value for the candidate models is given by:

$$\text{MSPE} : \left\{ m \in M : \frac{1}{N_\ell} \sum_{\ell=1}^{N_\ell} \left(y_{t+\ell} - \hat{y}_t^{(m)}(\ell) \right)^2 \right\}$$

```
N_L = length(nino_test)
# make predictions for held-out test data
ar1_forecast = predict(ar1, n.ahead = N_L)
ar6_forecast = predict(ar6, n.ahead = N_L)
mods = c("AR(1)", "AR(6)")
# compute MSPE for each
mspes = sapply(list(ar1_forecast, ar6_forecast),
  function(m){
```

```

        mean( (nino_test - m$pred)^2 )
    })
# report results
kable(t(round(mspes, 5)), col.names = mods,
      caption = "Observed MSPE of Candidates") %>%
  kable_styling(latex_options = "hold_position") %>%
  kable_styling(position = "center")

```

Table 4: Observed MSPE of Candidates

AR(1)	AR(6)
1.20036	1.01562

Then, to determine which method performs better, we select the argmin of the MSPEs set.

$$m_{\text{best}} = \underset{m \in M}{\operatorname{argmin}} \left\{ \text{MSPE}(m) \right\} = \underset{m \in M}{\operatorname{argmin}} \left\{ 1.200, 1.016 \right\} = \text{AR}(6)$$

In our case, this is the AR(6) model.

```
mods[which.min(mspes)]
```

```
## [1] "AR(6)"
```

Hence, by MSPE, the best model of the candidate models m_{best} remains the AR(6) model.

Part b.6.

On a single plot, display the test set of sea surface temperature, the predictions for both models, and their approximate 95% prediction intervals.

NOTE: We will include the plotting code for this question, as it specifically asks for it.

First, we compute the 95% prediction intervals using the following method: As before, we let $\hat{y}_t^{(m)}(\ell)$ be the $\ell \in \mathbb{Z}$ -step ahead Box-Jenkins forecast for model $m \in M$. We let $Z \sim N(0, 1)$ and “PrI” indicate the prediction interval at step ℓ , given significance level α and model m .

$$\text{PrI}(\ell \mid \{\alpha, m\}) = \hat{y}_t^{(m)}(\ell) \pm Z_{\alpha/2} \text{se}(\hat{y}_t^{(m)}(\ell))$$

For this question, we employ $\alpha = 0.05$ and $m \in M$ to compute the intervals. The R code to do this is below. We define the lower and upper intervals separately so that all of the values can be stored in a data frame for plotting.

```

alpha = 0.05
Z = qnorm(1-alpha/2)
# prediction intervals for AR(1)
ar1_Low = ar1_forecast$pred - Z*ar1_forecast$se
ar1_High = ar1_forecast$pred + Z*ar1_forecast$se
# prediction intervals for AR(6)
ar6_Low = ar6_forecast$pred - Z*ar6_forecast$se
ar6_High = ar6_forecast$pred + Z*ar6_forecast$se
# construct the data frame
bigplot_DF = data.frame(
  Time = as.Date(time(nino_test)),
  Observed = as.numeric(nino_test),
  Forecast_AR1 = as.numeric(ar1_forecast$pred),
  Lower_AR1 = as.numeric(ar1_Low),

```

```

Upper_AR1 = as.numeric(ar1_High),
Forecast_AR6 = as.numeric(ar6_forecast$pred),
Lower_AR6 = as.numeric(ar6_Low),
Upper_AR6 = as.numeric(ar6_High)
)

```

Now, we make the combined plot (with labels and a legend) below:

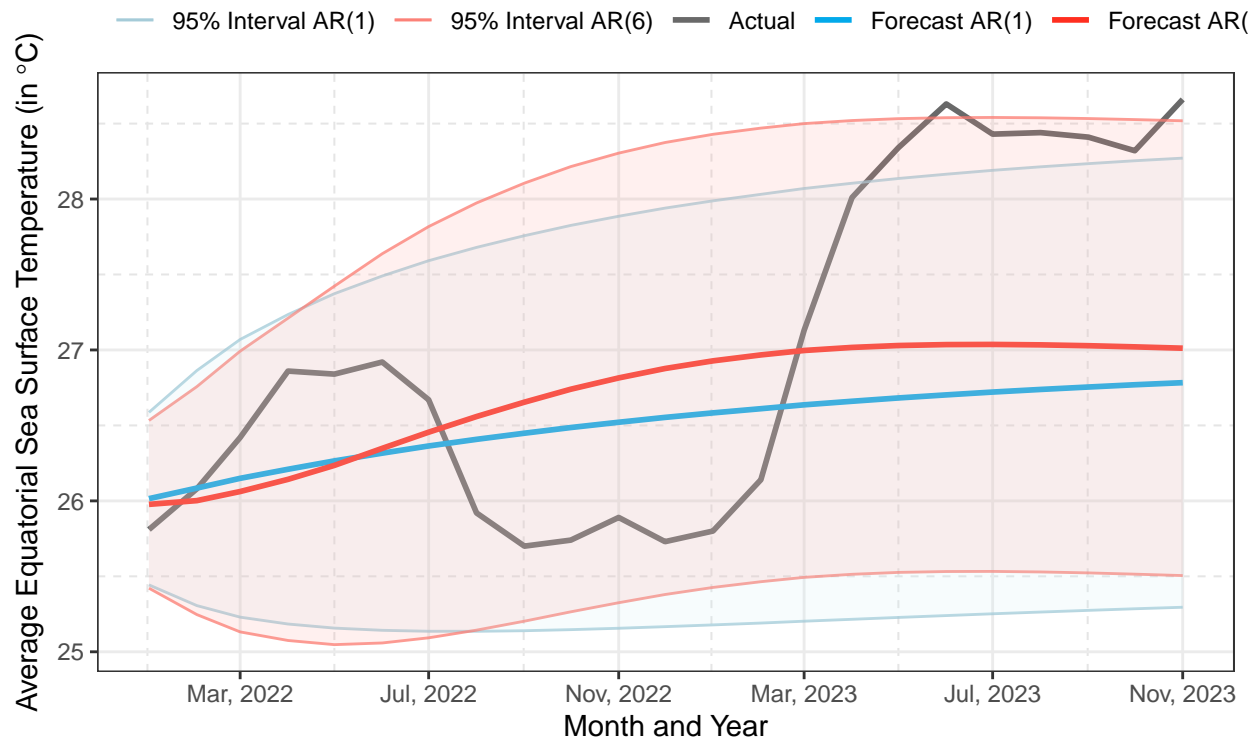
```

pfancy = ggplot(data = bigplot_DF, aes(x = Time)) +
  # lines for observed
  geom_line(aes(y = Observed, color = "Actual"), lwd = 1) +
  # lines and prediction interval for AR(1)
  geom_line(aes(y = Forecast_AR1, color = "Forecast AR(1)", lwd = 1) +
  geom_line(aes(y = Lower_AR1, color = "95% Interval AR(1)", alpha = 0.75) +
  geom_line(aes(y = Upper_AR1, color = "95% Interval AR(1)", alpha = 0.75) +
  # lines and prediction interval for AR(6)
  geom_line(aes(y = Forecast_AR6, color = "Forecast AR(6)", lwd = 1) +
  geom_line(aes(y = Lower_AR6, color = "95% Interval AR(6)", alpha = 0.75) +
  geom_line(aes(y = Upper_AR6, color = "95% Interval AR(6)", alpha = 0.75) +
  # light blue and red fill area between prediction intervals
  geom_ribbon(aes(ymin = Lower_AR1, ymax = Upper_AR1), fill = "#caf0f8", alpha = 0.15) +
  geom_ribbon(aes(ymin = Lower_AR6, ymax = Upper_AR6), fill = "#ff9991", alpha = 0.15) +
  # legend and colour assignment
  scale_color_manual(name = "",
    values = c(
      "Actual" = "#696969",
      "Forecast AR(1)" = "#02a9ea",
      "95% Interval AR(1)" = "#9fc8d6",
      "Forecast AR(6)" = "#ff2e1f",
      "95% Interval AR(6)" = "#fc796f"
    )
  ) +
  # customize x-axis for nice dates
  scale_x_date(date_breaks = "4 month", date_labels = "%b, %Y") +
  # add titles with units
  labs(
    title = "Box-Jenkins Forecasts and 95% Prediction Intervals",
    subtitle = "El Nino 3.4 Index from Jan. 2022 to Nov. 2023",
    x = "Month and Year",
    y = expression(
      paste("Average Equatorial Sea Surface Temperature (in ", degree, "C)")
    )
  ) +
  # customize theme: gridlines, margins, etc.
  theme_bw() +
  theme(
    panel.grid.minor = element_line(
      color = "grey90",
      linewidth = 0.35,
      linetype = "dashed"
    ),
    legend.position = "top",
    legend.justification = "left",
    legend.margin = margin(0,0,0,0)
  )
print(pfancy)

```


Box–Jenkins Forecasts and 95% Prediction Intervals

El Nino 3.4 Index from Jan. 2022 to Nov. 2023



Part c

Forecast sea surface temperature for 2022 and 2023 using the Holt-Winters method and the data from 1870-2021.

Part c.1.

Use the `HoltWinters` function in R to fit an appropriate model to the training data.

As in our original series (and with the Box-Jenkins method) we did not observe a significant trend, we will construct a Holt-Winters model with $\beta = 0$.

We will use an additive seasonal effect, which inherits the **Frequency** of the time series as the period. In this case, then, $p = 12$.

We then have an additive seasonal effect I_t for t given by the following:

$$I_t = \gamma(x_t - L_t) + (1 - \gamma)I_{t-p}$$

Where L_t is the level component that defines Holt-Winters smoothing/forecasting methods. Since $\beta = 0$, it is constructed without a T component, and is given by:

$$L_t = \alpha(x_t - I_{t-p}) + (1 - \alpha)L_{t-1}$$

For $\alpha, \gamma \in [0, 1]$. Further, for $\ell \in \mathbb{Z}$, the ℓ -step ahead forecast at time t is given by the following.

$$\hat{x}_t(\ell) = L_t + I_{t-p+\ell}$$

We report the coefficients from `HoltWinters` in the table below.

```

hw_no_trend = HoltWinters(nino_train, beta = FALSE, seasonal = "additive") # no trend
hw_results_coefs = (data.frame(
  Alpha = round(hw_no_trend$alpha, 3),
  Beta = paste(0, ".000", sep = ""),
  Gamma = paste(hw_no_trend$gamma, ".000", sep = "")
))
# format into table
rownames(hw_results_coefs) = NULL
kable(hw_results_coefs, "latex", escape = FALSE,
  col.names = c("$\\hat{\\alpha}$", "$\\hat{\\beta}$", "$\\hat{\\gamma}$"),
  caption = "Fitted Holt-Winters Coefficients") %>%
  kable_styling(latex_options = "hold_position") %>%
  kable_styling(position = "center")

```

Table 5: Fitted Holt-Winters Coefficients

$\hat{\alpha}$	$\hat{\beta}$	$\hat{\gamma}$
0.96	0.000	1.000

Our fitted $\hat{\gamma} = 1$, implying that the model places maximum smoothing weight on the current state of the season (i.e. the $(1 - \gamma)I_{t-p}$ component has zero weight according to the fitting process.) Further, our fitted $\hat{\alpha} \approx 0.96$ - this means that the model is placing a lot of significance on the current level (and period) and much less on levels in the past when fitting.

Part c.2.

Use this model to predict sea surface temperature from Jan 2022 through Nov 2023.

```

# make predictions and format as data frame
hw_predictions = data.frame(
  predict(hw_no_trend, n.ahead = 23, prediction.interval = TRUE, level = 0.95))


# prepare the hw forecast data
p4data = data.frame(
  Time = as.Date(time(nino_test)),
  Observed = as.numeric(nino_test),
  Forecast = as.numeric(hw_predictions$fit),
  Lower = as.numeric(hw_predictions$lwr),
  Upper = as.numeric(hw_predictions$upr)
)

print(p4)

```

Forecast and 95% Prediction Interval of Test Data: Holt–Winters

El Nino 3.4 Index from Jan. 2022 to Nov. 2023



Part c.3.

Calculate the mean squared prediction error.

The mean squared prediction error for our $N_\ell = 23$ total ℓ step ahead forecast is given by:

$$\text{MSE}_{\text{pred}} = \frac{1}{N} \sum_{\text{Holdout}} (\text{Truth} - \text{Forecast})^2 = \frac{1}{N_\ell} \sum_{\ell=1}^{N_\ell} (x_{t+\ell} - \hat{x}_t(\ell))^2$$

```
N = length(nino_test)
# verify equality in set cardinalities
stopifnot(all.equal(N, length(hw_predictions$fit)))
# calculate ms prediction error
ms_pre_hw = sum ( (nino_test - hw_predictions$fit )^2 ) / N
ms_pre_hw
```

```
## [1] 1.972297
```

The observed Mean Square prediction error for the Holt-Withers forecast is approximately 1.972.

Part c.4.

How does it compare to the Box-Jenkins models above?

```
# TODO: Comparison code goes here...
```

Question 2 : Hours Worked Forecasting

In this question we will predict the time series of monthly average values of the usual hours worked across all industries in Canada for the period from January 1987 until December 2023, which was explored in Assignment 1, using the file `usual_hours_worked_ca.csv`.

We'll use the Box-Jenkins method and Holt-Winters method.

Part 1: Data Preparation

Read in the data and create a time-series object for the mean monthly working hours:

```
hDF = read.csv("usual_hours_worked_ca.csv")
hours_series = ts(hDF$Hours, start = c(1987, 1), frequency = 12)
```


Separate the data into train and test. The training dataset should include all observations up to and including December 2020; The test dataset should include all observations from January 2021 to December 2023

```
hours_train = window(hours_series, end = c(2020, 12))
hours_test = window(hours_series, start = c(2021, 1), end = c(2023, 12))
# verify split is valid
stopifnot(all.equal(length(hours_test)+length(hours_train),
                     length(na.remove(hours_series))))
```

Prepare and plot the training data.

```
p5data = fortify.zoo(hours_train)
```

```
print(p5)
```



Part 2: Box-Jenkins Method

In this part, we select and fit a $\text{SARIMA}(p, d, q) \times (P, D, Q)_s$ model and make forecasts using the fitted model.

Part A: Differencing

Difference the training set time series at lag 1. Let $\{X_t\}_{t \in \mathbb{Z}}$ be our original series. We define below $Y_t = \nabla^1 X_t = X_t - X_{t-1}$


```
hours_train_diff = diff(hours_train, lag = 1, differences = 1)
```

Plot the new time series and its correlogram.

New Series

```
p6data = fortify.zoo(hours_train_diff)
```

```
print(p6)
```




Correlogram

```
p7data = data.frame(  
  h = 0:36,  
  rh = acf(hours_train_diff, plot = FALSE, lag.max = 36)$acf  
)
```

```
print(p7)
```

Correlogram of Differenced Series of Monthly Average Work Hours Using Training Data for Canada, Jan. 1987 to Dec. 2020



Comment on what you observe.

While the single-iteration ($d = 1$) of sequential differencing at lag 1 appears to have helped to mitigate the trend in the original training data, as seen in the plot of the differenced training time series, there still appears to be a notable seasonal component in the data. We know that this is likely the case due to the sinusoidal component in the correlogram, which is not really decreasing as a function of lag as we would like to see. This periodicity is indicative of a lingering seasonal component that we still need to difference out or remove via other methods.

Part B: Seasonal Differencing

Apply seasonal differencing to remove seasonal variation. Since we have monthly data, $s = 12$ seems sensible.


```
hours_train_diff_s = diff(hours_train_diff, lag = 12, differences = 1)
```

Plot the resulting differenced time series along with its sample acf and pacf. pure AR only for PACF [Series](#)

```
p8data = fortify.zoo(hours_train_diff_s)
```

```
print(p8)
```

Sequential and Seasonal [12] Difference of Monthly Average Hours Worked In Canada, Across all Industries from January 1987 to December 2020




Sample ACF

```
p9data = data.frame(  
  h = 0:80,  
  rh = acf(hours_train_diff_s, plot = FALSE, lag.max = 80)$acf  
)
```

```
print(p9)
```

Correlogram of Sequential and Seasonal [12] Differences

Monthly Avg. Work Hours, Canada, Jan. 1987 to Dec. 2020




Sample PACF

```
p10data = data.frame(  
  h = 1:80,  
  rh = pacf(hours_train_diff_s, plot = FALSE, lag.max = 80)$acf  
)
```

```
print(p10)
```


Sample PACF of Sequential and Seasonal [12] Differences

Monthly Avg. Work Hours, Canada, Jan. 1987 to Dec. 2020



Comment on what you observe:

The seasonal differencing paired with the sequential differencing seems to have done a good job in making the series stationary. We see from the plotted data as well as the correlogram that there isn't evidence of a significant season or trend. Further, the visual complexity of both the ACF and the PACF hints at the fact that there may be a blended model (i.e. a mixture of AR and MA) underlying this process. We will discuss in the next sections exactly what the parameters of the model equation are in this case.

Part C: Difference Parameters.

Based on the results of Part II (a) and (b), specify the values of d , D , and s .

From Part II (a), we applied one sequential difference using `diff`, so we know $d = 1$. Then, in Part II (b), we applied one seasonal difference using `lag = 12`, hence $D = 1$ with $s = 12$.

PACF ONLY FOR PURE AR

TAKE SOME MA MODEL

Part D: SARIMA Parameters

Based on the plots in Part II (b), suggest possible values of justifying your choices.

NOTE: It is very difficult to accurately determine these coefficients precisely, but I will justify my decisions to the best of my ability.

Firstly, let's discuss the presence of an MA component. We henceforth assume stationarity.

In the ACF, we see a fairly large r_h value at $h = 12$, that significantly cuts off for $h > 12$. Unlike a pure MA plot where we would see an expect a very sharp cutoff, in a blended model there will still be other components at work after the MA component.

All of the lags after $h = 12$ for $h \in [0, 36] \subset \mathbb{Z}$ are noteably smaller, as the Sample ACF shows.

This would imply that either we have a sequential MA with $q = 12$ or a seasonal MA with $Q = 1$. By the prinicple of parsimony, we prefer the latter. Also, this conclusion would imply that it is unlikely that there is a p component.

To explain, we recall the right-hand side of a SARIMA model, assuming $Q = 1$.

$$\begin{aligned}\text{RHS} &= \Theta(B^s)\theta(B)Z_t \\ \text{RHS} &= (1 + \tilde{\beta}B^{12})\theta(B)Z_t \\ \text{RHS} &= \theta(B)Z_t + \theta(B)\tilde{\beta}Z_{t-12}\end{aligned}$$

This would help explain the decently large “spike” at lag 12 (also perhaps suggesting that $\tilde{\beta} < 0$.)

However, is there a sequential component $q > 0$? By the above, under the assumption that $Q = 1$, if $\theta(B)$ was of an order greater than zero, there would be additional components in the right-hand side that we’d expect to see reflected in the ACF.

The most likely visually is that $q = 1$. An explanation of why is facilitated by the equation below:

$$\begin{aligned}\text{RHS} &= (1 + \beta B)Z_t + (1 + \beta B)\tilde{\beta}Z_{t-12} \\ \text{RHS} &= Z_t + \beta Z_{t-1} + \tilde{\beta}Z_{t-12} + \beta\tilde{\beta}Z_{t-13}\end{aligned}$$

In other words, from our understanding of MA models, we would see large in magnitude r_1 and (slightly less large) r_{13} . In the sample ACF, with a spike at lag 1 less than $-2/\sqrt{n}$ and a shorter spike at lag 13, slightly above $2/\sqrt{n}$, however, it is difficult to discern whether or not these are caused by the AR components of the model.

I am more inclined to favour the AR components causing these observed r_h values than additional MA components due to the fact that there is a large autocorrelation at lag 11, which isn’t described by this sort of MA formulation. Notably, the fact that visually $|r_{11}| \approx |r_1|$ lends informal credence to the hypothesis that it is AR components rather than MA causing the r_1 and r_{13} observations. If there were a $q > 0$, I would expect it to have a notable magnitude outside of the general AR pattern. So, while it’s theoretically plausible that $q > 0$, by this reasoning we settle onto $q = 0$ in favour of a more robust AR components.

Now let’s consider whether or not there are AR components to the series. Visually, the slowly-decaying oscillating pattern is indicative of one or more AR components. Further, since the lag doesn’t cut off to white noise immediately after the MA components (we see it sometimes creep above the $\pm 2/\sqrt{n}$) this tells us there is very likely an AR component of some variety.

To figure out exactly what order these components are, we utilize the PACF, recalling the assumption of stationarity.

However, we first must recall the left-hand side of the SARIMA equation, so we have a better understanding of the PACF Values

$$\begin{aligned}\text{LHS} &= \Psi(B^s)\psi(B)W_t \\ \text{LHS} &= \Psi(B^s)\psi(B)(\nabla_1^{12}\nabla^1 X_t) \\ \text{LHS} &= \Psi(B^s)\psi(B)(X_t - X_{t-12} - X_{t-1} + X_{t-13})\end{aligned}$$

So, it’s a bit more difficult to directly examine the PACF for the largest value of $\hat{\alpha}_{kk}$ such that $\forall i > k, |\hat{\alpha}_{ii}| < 2\sqrt{n}$ as we normally would, since it isn’t direct to discern which coefficient in the AR decomposition it actually aligns to in the overall model (i.e. in the full seasonal/sequential differenced with MA.)

Then we find a logical combination of p and P , we can informally examine the PACF for what seems “sensible.” What immediately jumps out to me is the presence of a large $\hat{\alpha}_{kk}$ at $k = 12$. This is indicative of a seasonal

AR component, i.e. $P \geq 1$. Further, there is an immediate jump at $\hat{\alpha}_{11}$ as well as other “spikes” across the PACF at $\hat{\alpha}_{11,11}, \hat{\alpha}_{23,23}$, etc. Since these values $\{1, 11, 23, \dots\}$ are not evenly divisible by 12, we know that there must exist one or more sequential AR components, i.e. $p \geq 1$. For the sake of parsimony, my initial suggestion would be to choose $P = p = 1$ as a starting point.

Due to the complexity inherent in this particular mixed models, it is difficult to visually determine values of p, P, q , and Q , however, from the arguments above, $\{p, P, q, Q\} = \{0, 1, 1, 1\}$ seems plausible.

```
# n = length(hours_train_diff_s)
# # (p,d,q): the AR order, the degree of differencing, and the MA order.
# p9data$rh[p9data$h]
# all_lags = acf(hours_train_diff_s, plot = FALSE, lag.max = n)$acf
# twelves = (all_lags[seq(from = 0, to = n, by = 1) %% 11 == 0])
#
# plot(0:(length(twelves)-1), twelves, type="n",
#       xlab="Index", ylab="Value", main="Lags at Step 12")
# segments(x0=0:(length(twelves)-1), y0=0, x1=0:(length(twelves)-1), y1=twelves)
# abline(h = 0)
```

Part E: Iterative AIC

Now use the Akaike’s Information Criterion (AIC) to select the model based on the training dataset in Part I. Fix the values of p and P as your suggestions in Part II (d),

Consider $q \in [0, 5]$ and $Q \in [0, 5]$.

Select the values of q and Q according to the AIC values.

Fit the model you choose and print the values of the estimated parameters along with the AIC value for the model.

First, we declare our fixed p, P, s, d and D values, then iterate through the different 25 models that can be developed and report the AIC for each.

```
p = 0; P = 0; s = 12; d = 1; D = 1
# declare empty matrix
results = matrix(0, nrow = 6, ncol = 6)
# with legible names
rownames(results) = c("q=0", "q=1", "q=2", "q=3", "q=4", "q=5")
colnames(results) = c("Q=0", "Q=1", "Q=2", "Q=3", "Q=4", "Q=5")
# iterate through all candidates
for (q in 0:5){
  for (Q in 0:5){
    results[q+1, Q+1] = arima(hours_train,
                              order = c(p, d, q),
                              seasonal = c(P, D, Q))$aic
  }
}
```

The AIC results are summarized in the table below.

```
kable(round(results, 3), caption = "AIC Values for Varying (q, Q)")
```

According to this table, the best model corresponds to:

```
which(results == min(results), arr.ind = TRUE)
```

```
##      row col
## q=4    5   6
```

Table 6: AIC Values for Varying (q, Q)

	Q=0	Q=1	Q=2	Q=3	Q=4	Q=5
q=0	-464.461	-613.121	-611.440	-609.462	-612.427	-612.305
q=1	-487.115	-616.650	-614.907	-612.958	-616.029	-616.225
q=2	-489.810	-618.001	-616.272	-614.388	-617.448	-618.428
q=3	-488.932	-622.962	-621.387	-619.411	-623.477	-623.785
q=4	-487.747	-624.882	-623.135	-621.160	-625.432	-625.629
q=5	-488.378	-623.821	-621.895	-619.917	-624.506	-624.633

```
# arima(hours_train, order = c(1, 1, 0), seasonal = c(1, 1, 0))$aic
# arima(hours_train, order = c(1, 1, 0), seasonal = c(1, 1, 1))$aic
min(results)
```

```
## [1] -625.6292
```

Which corresponds to $q = 4$ and $Q = 5$. It has an AIC of -627.61

Hence, we will use a model of the form:

$$\text{SARIMA}(0, 1, 4) \times (0, 1, 5)_{12}$$

The full model is given below:

```
best_fit = arima(hours_train, order = c(0, 1, 1), seasonal = c(1, 1, 5))
#?auto.arima(hours_train)

# Series: hours_train
# ARIMA(1,0,0)(0,1,4)[12] with drift
# auto.arima(hours_train, trace = F, ic = "aic",
#             start.p = 1, start.P = 1, stepwise = F,
#             max.Q = 5, max.q = 5, max.d = 1, max.D = 1)

# predict(best_fit, prediction.interval = T, level = 0.95)
```