

Lab 6: STAT 443

Caden Hewlett

2024-02-16

Question 1

The dataset `TempPG.csv` includes minimum temperatures measured at Prince George, BC, from 1919 to 2008. Read the data into R using either `read.table()` or `read.csv()` commands.

The column labelled “Summer” contains the summer minimum temperatures.

Part 1:

Extract those data, and coerce them into a time series object.


```
df = read.csv("TempPG.csv")
summer = df$Summer

summerts = ts(summer, start = 1918,
               end = 2008, freq = 1)
```


Part 2

Plot the time series, the sample acf and pacf.


```
par(mfrow = c(3, 1), mai = c(0.4, 0.6, 0.4, 0.4))
plot.ts(
  summerts,
  main = "Min. Temperature (C) Prince George, BC., 1919-2008",
  cex.main = 1.0,
  xlab = "Year",
  ylab = "Min. Temp. (Celsius)"
)
acf(summerts, main = "ACF of Summer Series Data")
pacf(summerts, main = "PACF of Summer Series Data")
```



ACF of Summer Series Data



PACF of Summer Series Data



Part 3

Comment on what you observe in these plots.

We see a slow (exponential) decline in magnitude in the ACF, implying that an AR model might be applicable. In an inspection of the PACF, we see that the final $\hat{\alpha}_{kk}$ value for significant coefficients is $k = 2$, this implies that if we were to fit an AR component it would have $p = 2$. However, there might be an additional (positive) trend component in the data (as seen in the first plot), that is causing us to see other values in the ACF. So, without first attempting to de-trend and perhaps difference the series, any conclusions made here are rudimentary.

Part 4

What ARMA model would you fit?

I would fit an $\text{ARMA}(p = 2, q = 0)$ model based on this analysis, for the reasons explained in the comments above; namely, that there appears to be an autoregressive component of order 2, and no visible moving average component.

Question 2

Fit the ARMA model you proposed above using the `arima()` command. Note that in the output of the `arima` command, 'intercept' refers to the mean of the process, which we denote by μ in class

```
fitted = arima(summerts, order = c(2, 0, 0))
fitted

##
## Call:
## arima(x = summerts, order = c(2, 0, 0))
##
## Coefficients:
##          ar1      ar2  intercept
##          0.4026  0.3446    7.0669
## s.e.      0.1029  0.1041    0.3263
##
## sigma^2 estimated as 0.6672:  log likelihood = -111.07,  aic = 230.15
```

Write down your fitted model.

Our original model (in pure theory) would be:

$$(X_t - \mu) = \alpha_1(X_{t-1} - \mu) + \alpha_2(X_{t-2} - \mu) + Z_t$$

Where $Z_t \sim \text{WN}(0, \sigma^2)$.

We now would have:

$$(X_t - \hat{\mu}) = \hat{\alpha}_1(X_{t-1} - \mu) + \hat{\alpha}_2(X_{t-2} - \hat{\mu}) + Z_t, \text{ where } Z_t \sim \text{WN}(0, \hat{\sigma}^2)$$

From the results, we see $\hat{\alpha}_1 = 0.4026$, $\hat{\alpha}_2 = 0.3446$, $\hat{\mu} = 7.0669$ and $\hat{\sigma}^2 = 0.6672$. In other words,

$$(X_t - 7.0669) = 0.4026(X_{t-1} - 7.0669) + 0.3446(X_{t-2} - 7.0669) + Z_t$$

Where $Z_t \sim \text{WN}(0, 0.6672)$.

Question 3

Use the `confint()` command to find 95% confidence intervals for relevant parameters.

(NOTE: We adopt the assumption of non-simultaneous confidence intervals, and hence do not apply a Bonferroni correction.)

The following will generate the CIs for $\hat{\alpha}_1$, $\hat{\alpha}_2$ and $\hat{\mu}$.

```
confint(fitted, level = 0.95)
```

```
##              2.5 %    97.5 %  
## ar1         0.2008682 0.6043640  
## ar2         0.1404444 0.5486957  
## intercept   6.4273702 7.7063323
```

If we wished, we could manually construct the confidence interval for the $\hat{\sigma}^2$ estimate, by assuming it takes a χ^2 distribution on $n - k$ degrees of freedom, where $k = 4$ for our 4 estimated parameters (estimates of α_1, α_2, μ and σ^2 itself.)

Let $\nu = n - k$. Recall that under the assumptions of the variance taking a chi-square distribution,

$$\text{CI} = \left\{ \frac{\nu \hat{\sigma}^2}{\chi_{\nu; 1 - (\alpha/2)}^2}, \frac{\nu \hat{\sigma}^2}{\chi_{\nu; (\alpha/2)}^2} \right\}$$

In R, this is:

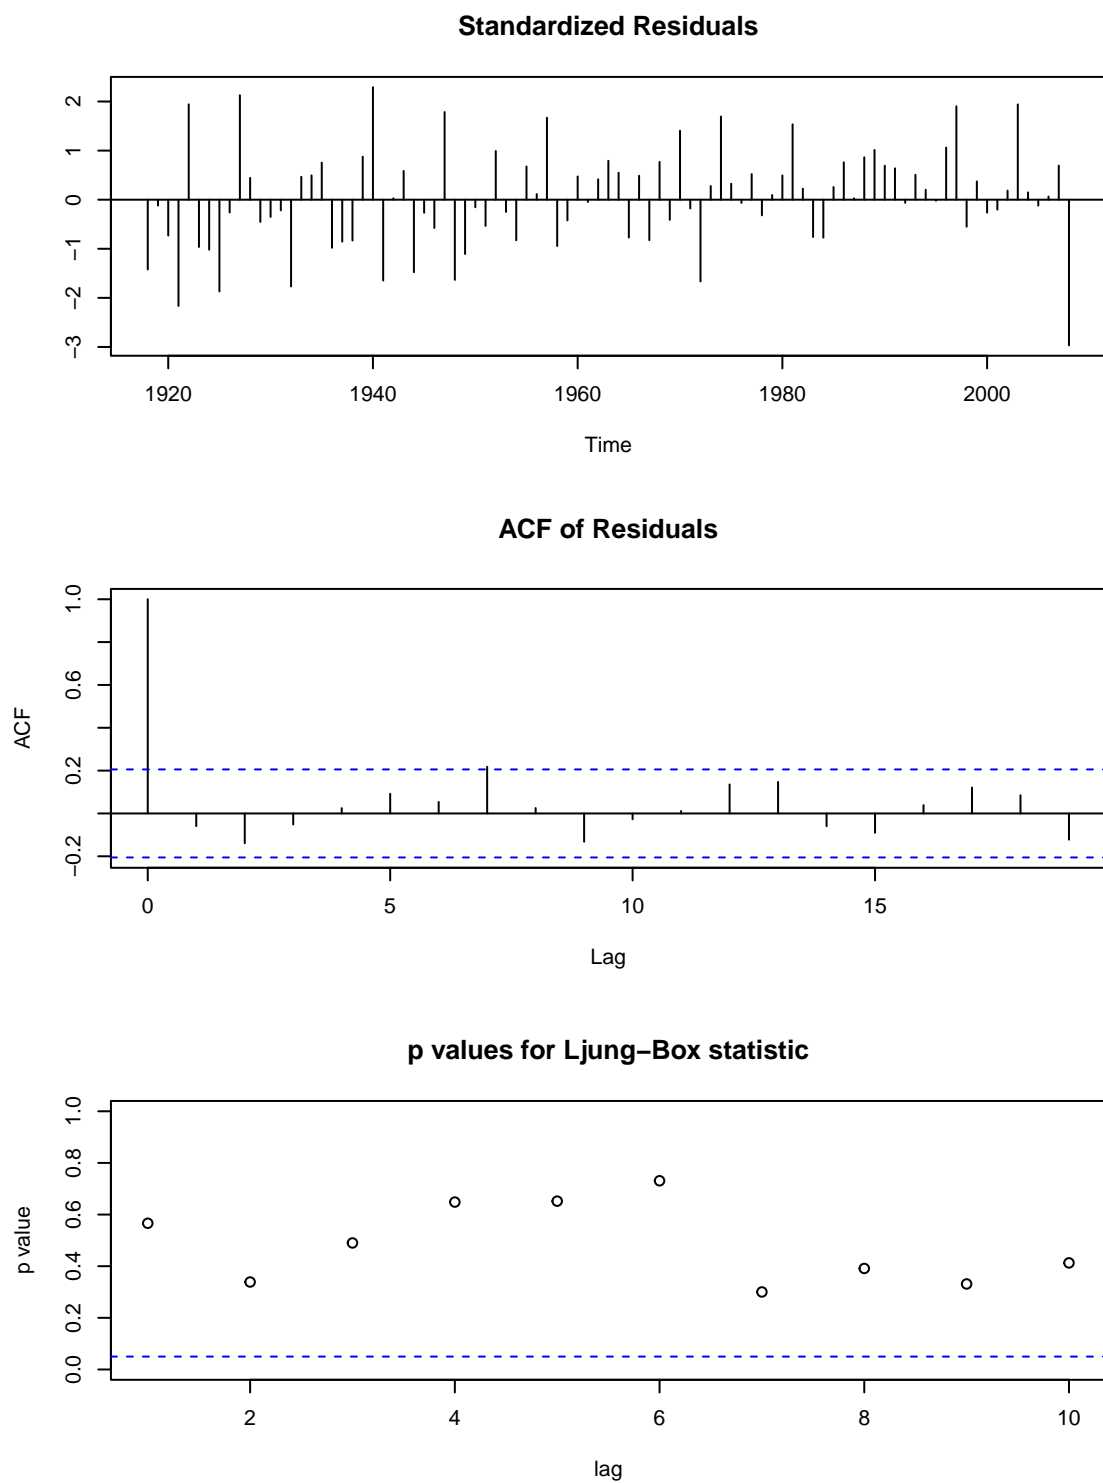
```
degf = length(summerts) - 4  
c((fitted$sigma2 * degf)/qchisq(0.975, degf),  
  (fitted$sigma2 * degf)/qchisq(0.025, degf))
```

```
## [1] 0.5061079 0.9200760
```

Question 4

Use the `tsdiag()` function to see diagnostic plots for the model you have fitted (remember to include “fig.height” option for a better display of your plots)

```
tsdiag(fitted)
```



Let's interpret these plots one-by-one

Plot 1

This plot shows a plot of the standardized residuals. We see that they are centered about zero (as we would anticipate if the fit was sound), and show no clear pattern - in a sense that it appears to simply be “noisy” values above and below zero. If there was a clear pattern (i.e. residuals consistently above/below), we would have less confidence in the strength of our fit; however, the lack of pattern in the residuals is encouraging.

Plot 2

Here, we see the autocorrelation function of the residuals. Beyond $\rho(0) = 1$, the majority of the values are within the $\pm 2/\sqrt{n}$ white noise boundary (save for lag 7 which *barely* exceeds it.) The fact that the residuals appear to be white noise reinforce our observations from the previous plot and demonstrate a good fit of the model.

Plot 3

The final plot shows the p -values of the Ljung-Box test statistic for different values of m .

We recall that this statistic takes a chi-square distribution where $Q_2 \sim \chi^2_{m-p-q}$, noting that we have $p = 2, q = 0$.

Crucially, there are no observed p -values below the $\alpha = 0.05$ threshold (horizontal line), implying that the lack-of-fit test would fail to reject H_0 ; in other words, at varying lag levels, there is not statistically significant evidence to suggest that the model does not in fact fit the data well. We see this in the Ljung-Box plot as the observed Q_2 values seem quite consistent with the corresponding χ^2 distribution on varying degrees of freedom (as we change m .)